

Technical Reference: Metaheuristics for the Hamiltonian Path Problem

Status: Living Document (Expanded)

1. Problem Definition & Graph Theory

1.1 The Hamiltonian Path

Formally, let $G = (V, E)$ be a graph where V is the set of vertices (nodes) and E is the set of edges. A **Hamiltonian Path** is a path that visits every vertex in V exactly once.

Unlike an Eulerian path (which visits every edge once and can be solved in polynomial time), determining if an arbitrary graph contains a Hamiltonian path is an **NP-Complete** problem. There is no known polynomial-time algorithm to solve it; as $|V|$ increases, the time required to solve it via brute force grows factorially ($O(N!)$).

1.2 The Environment: Erdős-Rényi Graphs

We generated our problem instances using the Erdős-Rényi model, denoted as $G(n, p)$.

- **N (Nodes):** The number of vertices (set to 50 in our main tests).
- **P (Probability):** The probability that an edge exists between any two distinct nodes (set to 0.3).
- **Significance:** Unlike grid graphs or lattices, Erdős-Rényi graphs are unstructured.

CRITICAL NOTE ON DIFFICULTY: A probability of $p = 0.3$ is significantly above the connectivity threshold for $N = 50$. The threshold for almost sure connectivity is $\frac{\ln N}{N} \approx 0.08$.

- **Implication:** At $p = 0.3$, the graph is dense, and Hamiltonian paths are abundant.
- **Recommendation:** To strictly test algorithmic robustness, future tests should lower p to the range of 0.10 – 0.15 to approach the "Phase Transition"—the point where the problem shifts from impossible to possible, which is where heuristics typically struggle.

1.3 Solution Representation (Encoding)

For all algorithms, a candidate solution is encoded as a **Permutation Vector**: $S = [v_1, v_2, \dots, v_n]$

- **Constraint:** Every integer from 0 to $N - 1$ must appear exactly once.
- **Validity Check:** In our metaheuristics, we enforce this constraint structurally (by using swap operators) rather than using penalties, ensuring we always search within the space of valid

permutations.

2. The Objective Function (Fitness Landscape)

We defined the "quality" of a path based on its adjacency continuity. The fitness landscape for this problem is discrete and highly "rugged," meaning there are many local optima (paths with only 1 or 2 breaks) that are difficult to escape.

2.1 Minimisation (Cost)

Used by Simulated Annealing and Tabu Search.

- Metric: "Broken Edges" (Gaps).
- Formula: $Cost(S) = \sum_{i=1}^{n-1} (1 \text{ if } (v_i, v_{i+1}) \notin E \text{ else } 0)$
- Goal: Reach a Cost of 0.
- Complexity: $O(N)$ per evaluation.

2.2 Maximisation (Fitness)

Used by the Genetic Algorithm.

- Metric: "Valid Edges".
- Formula: $Fitness(S) = \sum_{i=1}^{n-1} (1 \text{ if } (v_i, v_{i+1}) \in E \text{ else } 0)$
- Refinement (Squaring): To increase evolutionary pressure, we map this to $Fitness^2$. This widens the gap between a "good" solution (45 edges) and a "great" solution (49 edges), ensuring the elite solutions are selected more often.

3. Algorithm 1: Simulated Annealing (SA)

3.1 Concept

SA mimics the metallurgical process of annealing. In optimisation, "temperature" represents the probability of accepting a worse solution to escape local optima.

3.2 Key Mechanics

1. Neighbourhood Operator: Randomly swap two nodes in the permutation (2-exchange).
2. Acceptance Probability: If a new solution is better ($\Delta C < 0$), accept it. If it is worse ($\Delta C > 0$), accept it with probability: $P(accept) = e^{-\frac{\Delta C}{T}}$ Where T is the current temperature.
3. Cooling Schedule: Geometric Cooling: $T_{k+1} = \alpha \times T_k$ (Commonly $\alpha = 0.99$).
4. Behaviour: At high T , the algorithm behaves like a Random Walk (exploring). At low T , it behaves like Hill Climbing (exploiting).

4. Algorithm 2: Tabu Search (TS)

4.1 Concept

Tabu Search is a "Hill Climbing" algorithm with short-term memory. It maintains a list of "forbidden" (Tabu) moves to prevent the algorithm from cycling back to previously visited solutions.

4.2 Key Mechanics

1. **Stochastic Neighbourhood:** Instead of checking *all* possible swaps ($N(N - 1)/2$), which is computationally expensive ($O(N^2)$), we sample a subset (e.g., 50 random swaps). This effectively performs a "Stochastic Hill Climb".
2. **Tabu Tenure:** The number of iterations a move remains forbidden.
 - *Observation:* A tenure of 20-30 steps was necessary for $N = 50$.
3. **Aspiration Criteria:** A "safety valve" rule. If a move is on the Tabu list but results in a fitness better than the global best found so far, the Tabu status is ignored.

5. Algorithm 3: Genetic Algorithms (GA)

5.1 Concept

GA mimics Darwinian natural selection. A population of solutions evolves over generations via selection, crossover (breeding), and mutation.

5.2 The Permutation Challenge

Standard operators fail on permutations. We implemented specific operators to maintain valid paths:

1. **Ordered Crossover (OX1):**
 - Select a subsection from Parent A.
 - Copy it directly to the Child.
 - Fill remaining spots using the sequence from Parent B, *omitting* values already in the Child.
 - *Result:* Preserves relative ordering (schema) without duplicates.
2. **Swap Mutation:**
 - Swaps positions of two random values. This maintains the set of nodes but alters adjacency.
3. **Selection Mechanism (Missing in original notes):**
 - We utilise Tournament Selection (Size $k = 3$ or 5).

- *Reasoning:* Tournament selection allows for adjustable selection pressure (via k) and does not require sorting the whole population like Rank selection, making it computationally efficient.

4. Diversity Management:

- We tracked **Best Fitness** vs. **Mean Fitness**. If Mean Fitness approaches Best Fitness too quickly, the population has "collapsed" (low diversity).

6. Verification & Validation

6.1 Brute Force (Ground Truth)

- **Methodology:** Systematically generating every valid permutation ($N!$) and checking edge validity.
- **Limitation:** Feasible only for $N \leq 12$. For $N = 50$, the search space is approx 3×10^{64} , necessitating metaheuristics.

6.2 Stochasticity & Multi-Start

- **Issue:** Metaheuristics are non-deterministic.
- **Solution: Multi-Start Wrapper.** We run the algorithm k times (e.g., 5) and report the best result.
- **Success Metric:** We define "Success Rate" as the percentage of runs that achieve a Cost of 0 (perfect Hamiltonian Path).

7. Visualisation Strategy

Circular Layout (Ordered)

Standard force-directed layouts result in a "hairball" for random graphs. To prove a Hamiltonian Path was found:

1. Map the **Solution Path** to a perfect circle (Node 1 at 0° , Node 2 at x° , etc.).
2. Draw path edges in **Red** (perimeter).
3. Draw all other edges in **Faint Grey** (chords). This provides immediate visual verification: if the perimeter is unbroken red, the solution is valid.

8. Experimental Setup

For the sake of reproducibility, the following parameters define the "Base Case" for our comparison:

Parameter	Simulated Annealing	Tabu Search	Genetic Algorithm
Initialisation	Random Permutation	Random Permutation	Random Population
Stopping Condition	Cost = 0 OR Max Iter	Cost = 0 OR Max Iter	Fitness = Max OR Max Gen
Max Iterations	10,000	2,000	500 (Generations)
Specific Params	$T_{start} = 100, \alpha = 0.99$	Tenure=25, Neighbours=50	Pop=100, MutRate=5%
Runtime Complexity	Low (Single state)	Medium (Neighbourhood scan)	High (Population management)

9. Critical Analysis & Future Work

- **Phase Transitions:** As noted in Section 1.2, $p = 0.3$ is a dense graph. Future work should investigate performance near the phase transition ($p \approx 0.1$) where valid paths are rare and local optima are deeper.
- **Hybridisation:** The GA could be improved by adding a **Memetic** component—running a short Local Search (like 2-opt) on the best individuals of every generation to fine-tune the path.

10. Expected Comparative Results & Hypotheses

Based on the "No Free Lunch" theorem and the specific characteristics of the Hamiltonian Path problem (permutation-based, rugged landscape), we project the following outcomes:

10.1 Convergence Speed (Iteration Count)

Hypothesis: Tabu Search (TS) will require the **fewest iterations** to find a solution.

- **Reasoning:** TS employs an aggressive "best-improvement" strategy within its stochastic neighbourhood. Unlike SA, which "wanders" (random walk) at high temperatures, or GA, which relies on indirect recombination, TS actively selects the direction of steepest descent at every step.
- **Mechanism:** By evaluating 50 neighbours and picking the best one (even if it's non-improving, provided it's not Tabu), TS extracts more information per step than SA.

10.2 Computational Cost (Wall-Clock Time)

Hypothesis: Simulated Annealing (SA) will likely be the **fastest in real time**.

- **Reasoning:** While TS uses fewer iterations, each iteration is computationally expensive ($O(k \cdot N)$ where k is neighbourhood size). SA performs a single swap and evaluation ($O(1)$)

delta evaluation) per iteration. For easy instances (like $p = 0.3$), the low overhead of SA typically outweighs the "intelligence" of TS.

10.3 Solution Quality & Reliability

Hypothesis: For dense graphs ($p = 0.3$), all algorithms should achieve a near 100% success rate. However, as p decreases towards 0.1:

- **Tabu Search:** Likely to remain most robust. Its explicit memory prevents it from getting trapped in the deep local optima common in sparse graphs.
- **Genetic Algorithm:** Expected to perform worst in its pure form. Without a memetic component (local search), crossover operators often destroy the fragile adjacency links required for a Hamiltonian path, leading to stagnation near the solution (e.g., 2-3 broken edges).

11. Results & Graph Density Analysis

To rigorously test our hypotheses, we conducted an experiment on a significantly sparser graph ($p = 0.1$). This moves the problem closer to the critical phase transition, drastically increasing difficulty compared to the initial $p = 0.3$ tests.

11.1 Experimental Results ($N = 50, p = 0.1$)

The following data was extracted from our experimental runs (Source: `christos_mitsakopoulos.ipynb`).

Algorithm	Iterations / Gen	Result (Broken Edges)	Outcome
Simulated Annealing	3,000	11 Broken Edges	Failure
Tabu Search	2,000	8 Broken Edges	Failure
Genetic Algorithm	2,000	7 Broken Edges	Failure

Observation: Unlike the dense graph ($p = 0.3$) where solutions were trivial to find, **all three algorithms failed to find a valid Hamiltonian path at $p = 0.1$, converging to local optima with 7-11 gaps.**

11.2 Theoretical Analysis: The Phase Transition

The failure of the algorithms is not random; it is predicted by the properties of random graphs near the percolation threshold.

1. Connectivity Threshold

For an Erdős-Rényi graph $G(n, p)$ to be connected almost surely, p must satisfy:

$$p > \frac{\ln N}{N}$$

For $N = 50$:

$$p_{conn} \approx \frac{3.91}{50} \approx 0.078$$

At $p = 0.1$, we are very close to this limit. While the graph is likely connected, it is "barely" connected, meaning it relies on critical "bridges" that are hard to find.

2. The Hamiltonian Threshold

A stronger condition is required for a Hamiltonian Cycle. The celebrated result by Komlós and Szemerédi (1983) states that the critical probability p_{ham} is:

$$p_{ham} \approx \frac{\ln N + \ln \ln N}{N}$$

Substituting $N = 50$:

$$p_{ham} \approx \frac{3.91 + 1.36}{50} \approx \frac{5.27}{50} \approx 0.105$$

Conclusion: Our experimental $p = 0.1$ is actually **below** the theoretical threshold (≈ 0.105) for the almost-sure existence of a Hamiltonian cycle.

- **Result:** It is statistically probable that **no Hamiltonian path exists** in this specific graph instance.
- **Algorithmic Implication:** The algorithms did not fail due to poor design; they correctly identified that the best possible paths still contained gaps (7-11 broken edges). They were searching for a solution that mathematically did not exist.

11.3 Recommendations for Further Testing

To properly evaluate algorithmic performance without setting impossible tasks:

1. **Adjust p :** Increase p slightly to 0.15. This is safely above the Komlós-Szemerédi threshold ($0.15 > 0.105$), guaranteeing a solution exists while still being sparse enough to challenge the heuristics.
2. **Increase Iterations:** For sparse graphs, the search space contains deep local optima. SA should be run for 100,000+ iterations (up from 3,000) with a slower cooling schedule ($\alpha = 0.999$) to allow it to traverse these deep valleys.

12. Theoretical Benchmarks & The Komlós-Szemerédi Threshold

To elevate the analytical depth of the project, we must evaluate our results against established theoretical benchmarks rather than arbitrary pass/fail criteria.

12.1 The Komlós-Szemerédi Threshold (KS Limit)

The **Komlós-Szemerédi Theorem (1983)** provides the definitive "Phase Transition" point for Hamiltonicity in random graphs. It states that for a graph $G(n, p)$, if:

$$p(n) \geq \frac{\ln n + \ln \ln n + c_n}{n} \quad (\text{where } c_n \rightarrow \infty)$$

Then the graph almost surely contains a Hamiltonian cycle.

- **Significance for Metaheuristics:** This formula is the "Gold Standard" baseline. If an algorithm fails to find a solution when $p > p_{KS}$, the algorithm is flawed. If it fails when $p < p_{KS}$, the graph is likely impossible, and the algorithm is blameless.

12.2 Practical Benchmarks for Experimental Validation

To rigorously prove the effectiveness of the implemented algorithms, future experiments should consult the following benchmarks:

A. The "Phase Transition Plot" (Benchmark Standard)

Instead of testing single points ($p = 0.1$ or $p = 0.3$), the standard practice in combinatorial optimisation research is to generate a Phase Transition Plot:

- X-Axis: Edge Probability p (ranging from 0.05 to 0.20 in steps of 0.01).
- Y-Axis: Success Probability ($P(\text{Cost} = 0)$ over 50 runs).
- **Expectation:** You should observe a sigmoid (S-curve). The curve should shift from 0% to 100% sharply around $p \approx 0.11$ for $N = 50$.
 - *If your curve shifts right (e.g., success only starts at $p = 0.15$):* Your algorithm is inefficient.
 - *If your curve matches the theoretical threshold:* Your algorithm is optimal.

B. The Minimum Degree Check (Posá's Condition Wrapper)

Before running any expensive metaheuristic, a robust system should check Posá's necessary condition or simply the minimum degree $\delta(G)$.

- **Rule:** If $\delta(G) < 2$, a Hamiltonian Cycle is impossible. If fewer than 2 nodes have degree 1, a Hamiltonian Path is impossible.
- **Benchmark Utility:** This allows us to filter out "impossible" instances from the dataset,

ensuring that the "Failure Rate" metric reflects algorithmic failure, not graph impossibility.

C. Iterations at the "Edge of Chaos"

The most valuable performance data comes from the "Edge of Chaos"—the region just slightly above the KS threshold ($p \approx 0.12$ for $N = 50$).

- **Why:** In this region, paths exist but are extremely rare (narrow canyons in the fitness landscape).
- **Metric:** Compare the Number of Iterations to Convergence for SA vs. TS vs. GA specifically in this region ($0.11 < p < 0.12$). This distinction is important for understanding the behavior of the algorithms.