

Metaheuristic approach to the Hamiltonian Path

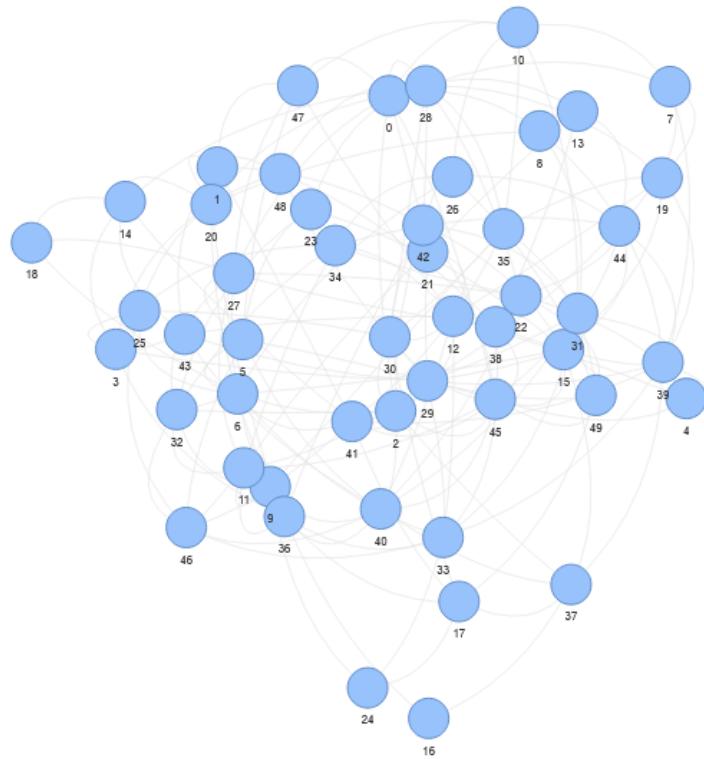


Figure 1: Visualised with Pyvis (Python), depicts the Hamiltonian path solved by the GA using the 2-OPT (inversion) operator & adaptive mutation, in the Erdős-Rényi random graph.

Module: Algorithms and Combinatorial Thinking 2025-2026

Written by:
Christos Christophoros Mitsakopoulos

Supervised by: Dr. Franck Delaplace

January 23, 2026

Contents

| | | |
|----------|---|-----------|
| 1 | Formal Definition and Introduction | 2 |
| 2 | Random Generation of the Problem: Erdős-Rényi Random Graph | 2 |
| 3 | Test Environment: Experimental Choices and Reasoning | 3 |
| 4 | Metaheuristic Algorithms: Mode of Action and Code Snippets | 4 |
| 4.1 | Objective Function | 4 |
| 4.2 | Simulated Annealing (SA) | 4 |
| 4.3 | Tabu Search (TS) | 5 |
| 4.4 | Genetic Algorithm (GA) | 5 |
| 4.4.1 | Adaptive Mutation Rate | 6 |
| 5 | Experiment: Comparing the Metaheuristic Algorithms | 6 |
| 5.1 | Batch Testing | 6 |
| 5.2 | Phase Transition | 7 |
| 6 | GA: Comparing Swap Operator to 2-Opt | 8 |
| 7 | Statistical Validation | 9 |
| 8 | Conclusion | 10 |
| 9 | Annex | 11 |

1 Formal Definition and Introduction

This report evaluates three metaheuristic approaches: *Simulated Annealing*, *Tabu Search*, and a *Genetic Algorithm*.

To start with, the Hamiltonian Path problem asks if a graph $G = (V, E)$ contains a path that visits every vertex / node $v \in V$ exactly once.

Given that the Hamiltonian Path problem is NP-Complete, as the number of vertices of the vertex set $n = |V|$ increases, the computational complexity required to solve the problem via brute force grows factorially ($O(n!)$). Consequently, exact algorithms become computationally intractable for large n , necessitating the use of metaheuristics – like the ones tested in this report. For this exact reason no evaluation of a brute force approach has been considered in the experiments of this report, as the computational and time complexity required is too high to practically implement.

In the context of bioinformatics, specifically *de novo* genome assembly, this problem is interesting. Arranging DNA reads in an acyclic graph can be modelled as finding a Hamiltonian path that maximises read overlap to create a contiguous sequence. Since this is a bioinformatics course, a genetic algorithm appropriated from the course material (much like all other algorithms implemented) was tested to match the theme of genome assembly heuristics.

Find all relevant results / as well as reproduce the experiment using the `main.py` at the following address: https://github.com/cmitsakopoulos/Delaplace_coursework. The Python script automates all tests demonstrated in this report and moreover, accepts user arguments through the command line interface; intended for tweaking parameters regarding base graph generation. The `-mode` argument allows running specific experiments:

```
1 # Default: runs all standard experiments
2 python main.py --nodes 50 --prob 0.15
3
4 # Specific modes:
5 python main.py --mode batch      # Batch testing only
6 python main.py --mode stats     # Statistical tests (Kruskal-Wallis,
    Wilcoxon)
7 python main.py --mode adaptive # Compare standard vs adaptive mutation
    GA
```

Listing 1: Available experiment modes. The default mode "all" runs batch, optimisation, phase transition, and visualisation experiments.

2 Random Generation of the Problem: Erdős-Rényi Random Graph

To benchmark the metaheuristic algorithms, an Erdős-Rényi (ER) $G(n, p)$ model was used. In which:

- n : The number of vertices in the graph.
- p : The probability that an edge exists between any two distinct vertices.

A larger parameter p will by effect increase the likelihood of finding a Hamiltonian Path, given that the probability of any two nodes having a connecting edge is larger. This was evident when first trying out the `main.py` Python script, where preliminary tests with a p of ≈ 0.3 and $n = 50$ showed that a Hamiltonian Path was indeed mathematically possible to find. It was so much so that all algorithms would converge to a zero cost solution (no broken edges).

```
1 # n_default = 50, p_default = 0.1, seed hardcoded to 42
2 g_exp = nx.erdos_renyi_graph(n=num_nodes, p=prob, seed=SEED)
```

Listing 2: Small snippet of code in which NetworkX in Python can generate an ER graph of chosen parameters, identify that the default options are intended for a "challenging" benchmark for the algorithms. Hardcoding the seed was not a deliberate analytical choice, just to ensure reproducibility of the graph itself – the algorithms are stochastic too...

3 Test Environment: Experimental Choices and Reasoning

To ensure significance of the results, each metaheuristic algorithm was executed for $N = 30$ runs. Relying on the Central Limit Theorem, the sampling distribution of the mean approximates a normal distribution as $N \geq 30$, even if the underlying population distribution is non-Gaussian. While increasing N reduces the Standard Error of the Mean (SEM), the precision improves only with the square root of N (i.e., $SEM = \sigma/\sqrt{N}$).

Additionally, in order to observe the impact on phase transition across all metaheuristics tested, the Komlós–Szemerédi theorem/threshold was used to compute the theoretical p -limit of the problem graph, specifically using the following equation:

$$p_{\text{limit}} = \frac{\ln(n) + \ln(\ln(n))}{n} \quad (1)$$

While $p < p_{\text{limit}}$, the probability of a Hamiltonian path existing approaches zero; thus any algorithm claiming a valid solution is likely "hallucinating", being overconfident (due to verification errors), or has encountered a statistical anomaly (highly unlikely in my testing case). Conversely, for $p > p_{\text{limit}}$, a path almost surely exists: consequently, if an algorithm fails to converge to a zero-cost solution in this region (zero broken edges), it shows that the algorithm is underperforming rather than the problem being impossible.

Another important consideration, was to examine the impact of switching from a "Swap" operator to an **"Inversion" (2-OPT) operator** between testing cases. This was also tested with the GA to demonstrate how 2-OPT can be a positive addition to an already accurate algorithm. Additionally, it was tested with SA and TS to demonstrate that the impact of this change is not algorithm-specific. This phenomenon should become evident when looking at the **"Phase Transition"** experiment of this report, where the 2-OPT operator was used instead of the Swap operator.

To move beyond visual comparisons and quantify whether observed performance differences are statistically meaningful, non-parametric statistical tests were employed. The Kruskal-Wallis H-test determines whether at least one algorithm performs significantly

differently from the others, while pairwise Wilcoxon signed-rank tests identify which specific pairs differ. These non-parametric tests are appropriate because solution quality scores do not follow a normal distribution—algorithms either find optimal solutions (cost = 0) or get stuck at various local optima. Additionally, Cohen’s d effect size quantifies the practical magnitude of differences, distinguishing between statistically significant results and practically meaningful ones. An effect size $|d| \geq 0.8$ indicates a large, practically important difference between algorithms.

4 Metaheuristic Algorithms: Mode of Action and Code Snippets

4.1 Objective Function

The aim of these algorithms is to identify a permutation S of vertices that minimises the number of broken edges in the path. For a graph $G = (V, E)$ and a candidate path $S = [v_1, v_2, \dots, v_n]$, the cost function $C(S)$ is defined as:

$$C(S) = (n - 1) - \sum_{i=1}^{n-1} \mathbb{I}((v_i, v_{i+1}) \in E) \quad (2)$$

Where \mathbb{I} is an indicator function that equals 1 if the edge exists and 0 otherwise. A global optimum is reached when $C(S) = 0$ – a Hamiltonian Path.

4.2 Simulated Annealing (SA)

Simulated Annealing explores solutions by accepting both better or even worse solutions, based on a probability that decreases over time (referred to as parameter T). This logic prevents the algorithm from arriving at a final solution before reaching a truly optimal or near optimal solution. The probability P of accepting a new solution S' with cost difference $\Delta C = C(S') - C(S)$, is governed by the Metropolis criterion (see also the Python implementation below):

$$P(\text{accept}) = \begin{cases} 1 & \text{if } \Delta C < 0 \\ e^{-\frac{\Delta C}{T}} & \text{if } \Delta C \geq 0 \end{cases} \quad (3)$$

```

1 # From main.py: Calculate cost difference
2 delta = neighbor_cost - current_cost
3
4 # Accept if better (delta < 0) OR with probability exp(-delta/T)
5 if delta < 0 or random.random() < exp(-delta / temp):
6     current = neighbor
7     current_cost = neighbor_cost

```

Listing 3: *Metropolis Criterion: As the loop progresses, the temperature T decays geometrically ($T_{k+1} = 0.985 \cdot T_k$), gradually turning the search into a simple greedy descent (hill climbing)*

4.3 Tabu Search (TS)

Tabu Search differs from SA by using a deterministic, memory-based approach. TS explores the immediate neighbourhood of a current solution, then moves to the best available neighbouring solution, even if that neighbouring solution is worse than the current solution.

To prevent cycling (revisiting the same solutions again and again), the algorithm maintains a *Tabu List* – a short-term memory that prevents recent solutions for a specific duration, called tenure.

```
1 # From main.py: Moving to the best candidate not in the Tabu list
2 if cand not in tabu_list:
3     current = cand
4     found_move = True
5
6 # Update memory
7 tabu_list.append(current)
8 if len(tabu_list) > tenure:
9     tabu_list.pop(0) # Remove oldest entry
```

Listing 4: Tabu Search Memory Logic

4.4 Genetic Algorithm (GA)

The Genetic Algorithm attempts to mimic natural selection. Unlike SA and TS, which improve a single solution, GA evolves a population of solutions. The operator for permutation is *Ordered Crossover* (OX1), which is important because standard single-point crossover would result in duplicate or missing vertices / nodes.

The code uses OX1 to preserve the ordering of a sub-segment from one "parent" while filling the remaining slots with genes from the second "parent".

```
1 # From main.py: Preserves sub-segment from parent 1
2 child[start:end] = p1[start:end]
3
4 # Fills remaining slots from parent 2, skipping duplicates
5 for i in range(size):
6     if child[i] is None:
7         while p2[current_p2_idx] in child:
8             current_p2_idx += 1
9         child[i] = p2[current_p2_idx]
```

Listing 5: Ordered Crossover (OX1) Implementation

Using the standard `swap` operator (exchanging two indices) disrupts the adjacency of the path significantly. In comparison, the `inversion` (2-Opt) operator reverses a segment of the path. This is mathematically better for path search problems because it preserves the internal adjacency of the reversed segment, only breaking the two edges at the endpoints of the segment. This distinction is implemented via the ‘`op_inversion`’ function in ‘`main.py`’ and is the primary driver for convergence in denser graphs. A graph is produced at the end during testing to demonstrate the differences of both applications.

$$\text{Swap}(S) \rightarrow \text{High disruption of edges} \quad (4)$$

$$\text{Inversion}(S) \rightarrow \text{Minimal disruption (2-Opt)} \quad (5)$$

4.4.1 Adaptive Mutation Rate

Another improvement over that of the mutation operator (Swap vs 2-Opt), which determines how to mutate a solution in order to develop it, is improving on how often to mutate. With a high fixed mutation rate, the algorithm promotes exploration but disrupts good solutions, while low mutation rates enable exploitation but risk converging early; thereby leading to suboptimal solutions. An adaptive approach addresses this by adjusting the mutation probability μ based on population diversity D :

$$D = \frac{|\mathcal{F}_{\text{unique}}|}{P}, \quad \mu_{\text{adaptive}} = \mu_{\text{base}} \cdot (1 + (0.5 - D)) \quad (6)$$

When $D < 0.5$ (population becoming homogeneous), mutation increases to reintroduce variation. When $D > 0.5$ (healthy diversity), mutation decreases to exploit promising solutions. The rate is clamped to $[0.1, 0.6]$. This self-regulating mechanism helps prevent premature convergence without manual tuning.

```

1 def calculate_diversity(fitnesses):
2     return len(set(fitnesses)) / len(fitnesses)
3
4 # Inside GA loop, per generation:
5 if adaptive:
6     diversity = calculate_diversity(fitnesses)
7     current_mut = base_mut_rate * (1 + (0.5 - diversity))
8     current_mut = max(0.1, min(0.6, current_mut))

```

Listing 6: *Adaptive mutation rate logic from main.py. Diversity is computed per generation, and mutation rate is adjusted accordingly.*

For a fixed mutation rate μ_0 , the change in diversity that is expected per generation matches a drift-selection balance. When diversity is low ($D \rightarrow 0$), the population converges early; such that all individuals occupy similar regions of the search space. The adaptive scheme $\mu_{\text{adaptive}} = \mu_0 \cdot (1.5 - D)$ gives a self-correcting feedback loop: when $D < 0.5$, mutation increases to escape local optima; when $D > 0.5$, mutation decreases to enable for more promising solutions. The equilibrium $D^* = 0.5$ represents an optimal exploration-exploitation balance.

5 Experiment: Comparing the Metaheuristic Algorithms

5.1 Batch Testing

The experiment code in `main.py` has been hardcoded to repeat the stochastic run of each metaheuristic algorithm 30 times, with the user chosen parameters upon initialization through the command line. The results depicted in [Figure 2](#) and [Figure 3](#), were computed with an ER problem graph of $n = 50$ and $p = 0.15$.

Briefly, with reference to [Figure 2](#), the lowest solution quality (highest number of broken edges) is demonstrated by SA, then TS and with the GA, being the best performer. Inversely, the highest execution time per batch run, was demonstrated by GA, followed by TS and lastly SA. These results confirm the expected correlation between computational cost and solution quality (indicating the code works as intended). SA is the fastest because it performs a single $O(1)$ evaluation per step, but its single-trajectory stochasticity struggles to escape bad solutions (local optima) in a solution landscape. TS evaluates a

neighbourhood of size $k = 50$ per step ($O(k)$) (focus on local searches), which linearly increases runtime. Lastly, the GA achieves a near-optimal solution albeit, at the highest computational cost. It maintains diversity through a population-based search ($O(P \cdot N)$ per generation), allowing it to traverse the complex fitness landscape more effectively than the trajectory-based methods.

Nevertheless, the results in [Figure 2](#) clearly show that all algorithms – apart from the GA – are greatly underperforming in this testing case. For an ER graph with $n = 50$, the Komlós Szemerédi limit is $p = 0.106$, considering that the chosen – testing – $p = 0.15$, is greater than the theoretical limit, a solution should be mathematically infeasible; albeit, difficult to obtain.

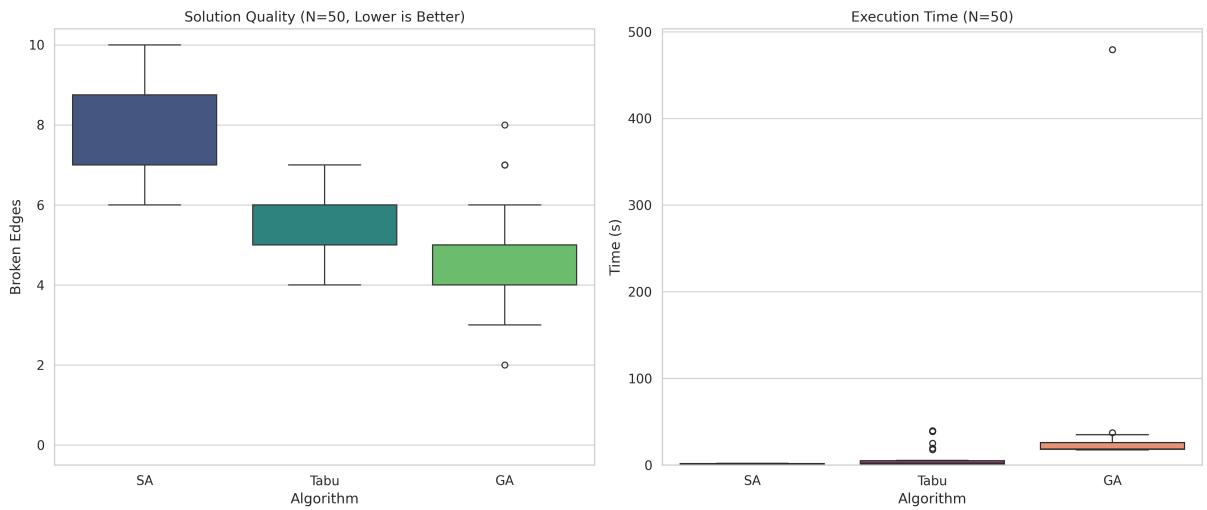


Figure 2: In the box whisker plots depicted, "Solution Quality" (Y-axis: Broken Edges) is shown on the left and "Execution Time" (Y-axis: Time(seconds)) on the right. As discussed prior, due to the low probability of edges between any two nodes, the graph is sparse and the metaheuristic algorithms – while performing differently from one another – are not successful in this testing case. As such, the amount of broken edges seen between algorithms is not unlikely, and the inverse relationship between the number of broken edges and execution time is a clear demonstration of algorithmic complexity – which leads to positive outcomes in problem solving.

5.2 Phase Transition

Briefly, with reference to [Figure 3](#), there appears to be overconfidence of the TS and GA algorithms, which demonstrate a non-zero level of success at a mathematically improbable rate of $p = 0.1$; where $p = 0.1 < p_{\text{limit}} = 0.106$. Given that the difference between p and p_{limit} is $\Delta p = 0.006$, this could be believed to be a statistical anomaly that occurred in this new 30-run test – compared to the previous batch testing scenario. After introducing a new `argparse` argument to the `main.py` file to exclusively re-run the phase transition experiment, the results – over 5 re-runs – led to the same conclusions as in [Figure 3](#). One could argue that due to a **switch from the "Swap" operator to the more optimised "Inversion" (2-Opt) operator**, the GA and TS algorithms are able to find a global optimum at a mathematically improbable rate of $p = 0.1$, as well as be considerably successful at higher graph densities ($p > 0.1$); all the while SA is equally performing better than before (than in [Figure 2](#)).

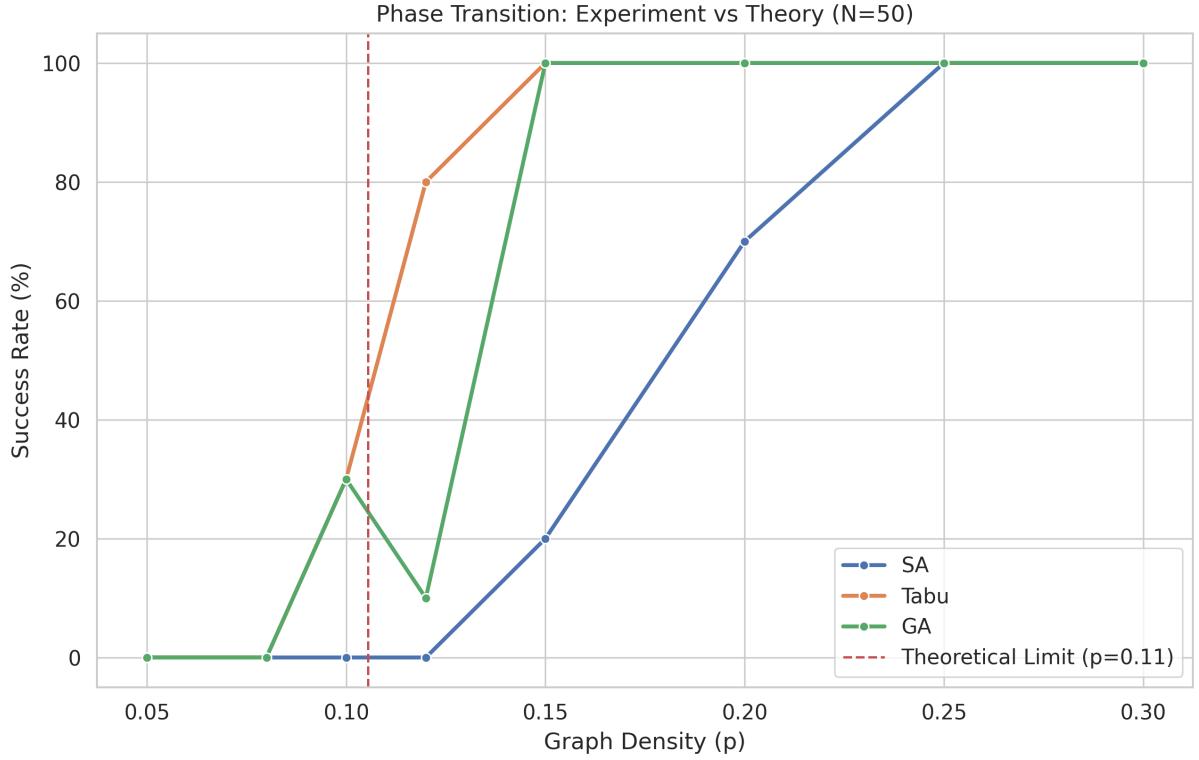


Figure 3: This line chart depicts the phase transition of the proposed problem Erdos-Renyi graph, at a cycling p rate between 0.05 and 0.3. The theoretical limit is computed automatically based on the input problem graph, given the Komlós–Szemerédi formula. Observe overconfidence from the TS and GA algorithms, which demonstrate success at a mathematically improbable rate of p . The success rate of all algorithms is congruent with previous results.

6 GA: Comparing Swap Operator to 2-Opt

Congruent with the phase transition results, it is evident that the switch from the Swap operator (Figure 2) to the Inversion (2-Opt) operator (Figure 3), improves the performance of the GA (see Figure 4). Importantly, looking at the number of broken edges at each iteration, it is clear that the 2-OPT based GA converges at a faster and smoother, or more consistent pace, than the Swap based GA. To briefly reiterate on what was said prior on the report, this phenomenon can be due to two complementary reasons:

Firstly, the 2-OPT (or inversion) operator is less destructive at each iteration. It reverses a segment between indices i and j at each iteration, removing two edges (v_i, v_{i+1}) and (v_j, v_{j+1}) , and replacing them with (v_i, v_j) and (v_{i+1}, v_{j+1}) , whilst preserving the adjacency of all nodes within the segment. As such, the 2-OPT operator enables the algorithm to make small changes which are beneficial at each iteration. Therefore, as can be seen in Figure 4, the GA with the 2-OPT operator improves its global solution at each iteration, whereas the Swap operator will repeatedly plateau/stagnate due to destructive solutions between iterations. Secondly, as mentioned before, the Swap operator exchanges two nodes v_i and v_j arbitrarily, which can affect up to four edges: (v_{i-1}, v_i) , (v_i, v_{i+1}) , (v_{j-1}, v_j) , and (v_j, v_{j+1}) . Instead of making minor improvements locally, the Swap operator will lead the algorithm to more destructive changes at each iteration. In turn, this will cause the search

to stagnate (maintain one solution over multiple iterations) over a certain amount of iterations. While a more beneficial solution can be found accidentally, the algorithm risks getting stuck in local minima; minima which can be avoided by the 2-OPT operator.

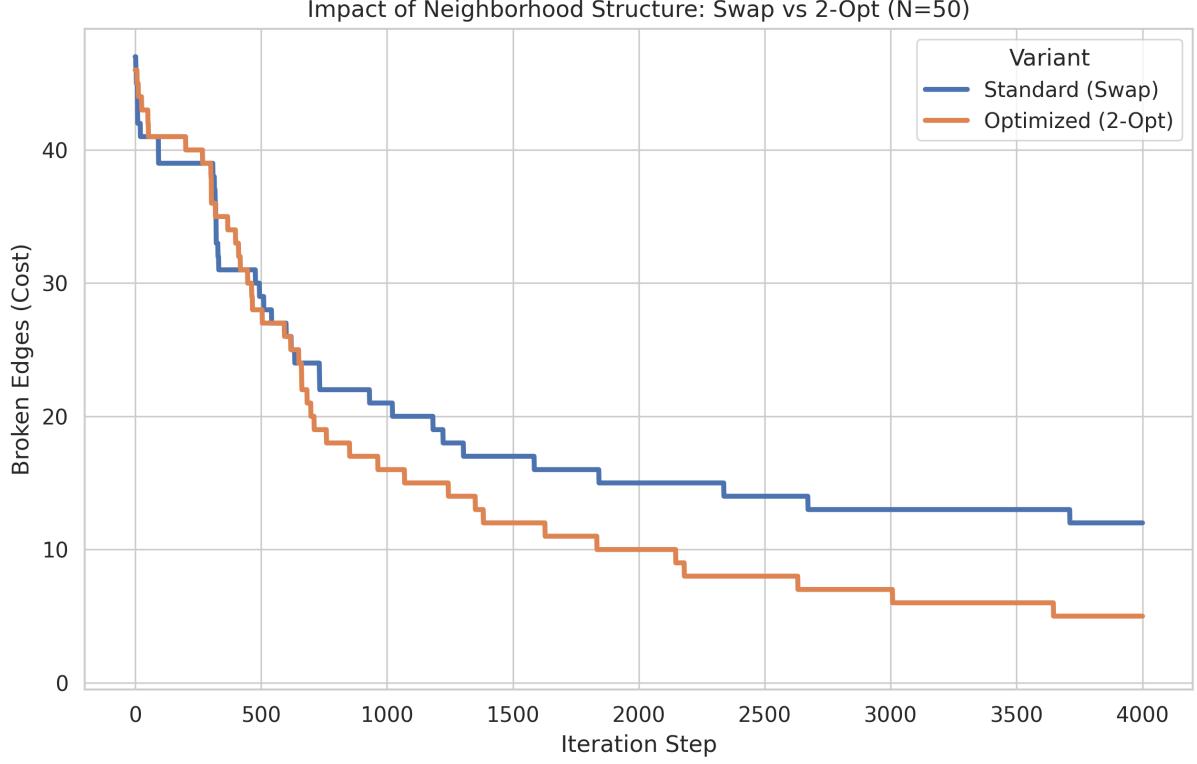


Figure 4: In the depicted line chart, Broken Edges (Y-axis) for which the objective function seeks to minimise (Cost) by maximising the number of edges in the path, is plotted against the number of generations – repeated runs (X-axis). The GA with the Swap operator is shown in blue, while the GA with the Inversion (2-Opt) operator is in orange. Observe the clear disparity between the two operators, in which it's clear that the 2-OPT operator GA converges much faster to a near-global optimum (zero broken edges) than its Swap-GA counterpart.

7 Statistical Validation

To quantify the statistical significance of the performance differences observed in the figures above, the Kruskal-Wallis H-test was applied:

$$H = \frac{12}{N(N+1)} \sum_{i=1}^k \frac{R_i^2}{n_i} - 3(N+1) \quad (7)$$

A significant result ($p < 0.05$) confirms that the algorithms perform differently. Pairwise Wilcoxon signed-rank tests further identify which specific pairs differ:

$$W = \min \left(\sum_{d_i > 0} R_i, \sum_{d_i < 0} R_i \right) \quad (8)$$

Finally, Cohen's d effect size quantifies the magnitude of these differences:

$$d = \frac{\bar{x}_1 - \bar{x}_2}{s_{\text{pooled}}} \quad \text{where} \quad s_{\text{pooled}} = \sqrt{\frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2}} \quad (9)$$

I visualise this statistical analysis through a three-panel figure containing:

- Box plot with Mann-Whitney significance annotations (star notation: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$).
- Mean broken edges with 95% confidence interval error bars.
- Effect size heatmap showing pairwise Cohen's d values.

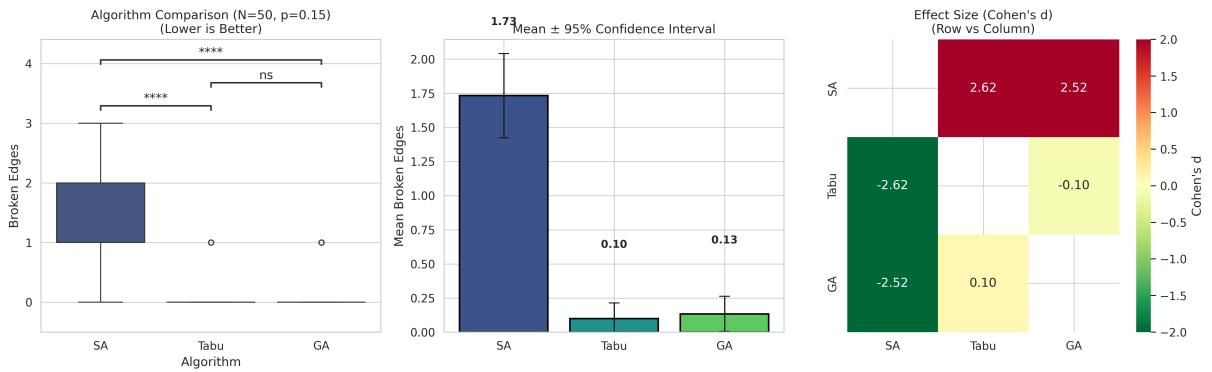


Figure 5: *Statistical comparison of SA, Tabu Search, and GA. Left: Box plot with Mann-Whitney significance annotations. Centre: Mean broken edges with 95% CI error bars. Right: Cohen's d effect size heatmap showing pairwise comparisons.*

With reference to Figure 5, SA exhibits a mean of $\bar{x}_{SA} = 1.73$ broken edges (95% CI: between 1.42 and 2.04), while Tabu Search ($\bar{x}_{Tabu} = 0.10$, CI: between -0.01 and 0.21) and GA ($\bar{x}_{GA} = 0.13$, CI: between 0.00 and 0.26) achieve near-optimal solutions. Importantly, but expectedly so, the Mann-Whitney tests confirm highly significant differences between SA and both other algorithms ($p < 0.0001$), whereas Tabu vs GA is non-significant. Equally, the effect size heatmap shows large differences between all: $d_{SA \leftrightarrow Tabu} = 2.62$ and $d_{SA \leftrightarrow GA} = 2.52$ (both $|d| > 0.8$), while $d_{Tabu \leftrightarrow GA} = -0.10$ indicates negligible difference. These results statistically confirm that the memory-based (Tabu) and population-based (GA) strategies substantially outperform the single-trajectory stochastic approach (SA) on this graph configuration. However, this strong distinction diminishes at higher graph densities ($p \geq 0.15$), where all algorithms achieve 100% success as shown in the phase transition analysis.

8 Conclusion

In this report I compare three metaheuristic approaches for the Hamiltonian Path problem on specific Erdős-Rényi random graph. Out of the three, the Genetic Algorithm and Tabu Search consistently outperformed Simulated Annealing, with statistically significant differences (Cohen's $d > 2.5$) at moderate graph densities. The addition of a 2-OPT (inversion) operator, proved better than the standard swap operator, thus leading to faster and smoother convergence. Expectedly, the phase transition analysis confirmed that all

algorithms achieve 100% success above the Komlós & Szemerédi threshold ($p > p_{\text{limit}}$), though GA and Tabu demonstrated surprising success even slightly below this theoretical limit; with GA excelling the most. For most applications, the GA with 2-OPT and adaptive mutation represents the best choice, albeit at higher computational cost.

9 Annex

```

1 import random
2 import time
3 import math
4 import itertools
5 import argparse
6 import numpy as np
7 import pandas as pd
8 import networkx as nx
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11 from math import exp, log
12 from scipy import stats
13 from itertools import combinations
14
15 # Try to import statannotations for publication-ready p-value
16 # annotations
16 try:
17     from statannotations.Annotator import Annotator
18     HAS_STATANNOTATIONS = True
19 except ImportError:
20     HAS_STATANNOTATIONS = False
21     print("Note: Install 'statannotations' for p-value annotations on
22 plots: pip install statannotations")
22
23 # --- CONFIGURATION ---
24 SNS_THEME = "whitegrid"
25 SNS_CONTEXT = "paper"
26 SEED = 42
27
28 # Apply settings
29 random.seed(SEED)
30 np.random.seed(SEED)
31 sns.set_theme(style=SNS_THEME, context=SNS_CONTEXT, font_scale=1.2)
32
33 # =====
34 # 1. CORE HELPER & FITNESS FUNCTIONS
35 # =====
36
37 def get_initial_solution(nodes):
38     """Returns a random permutation of nodes."""
39     p = list(nodes).copy()
40     random.shuffle(p)
41     return p
42
43 def count_edges(path, graph):
44     """Counts valid edges in the path for a specific graph."""
45     edges = 0
46     # Optimization: Iterate via index to avoid creating new lists
47     for i in range(len(path) - 1):

```

```

48         if graph.has_edge(path[i], path[i+1]):
49             edges += 1
50     return edges
51
52 def fitness_min(path, graph):
53     """Minimisation: Returns number of broken edges. Target = 0."""
54     target = len(path) - 1
55     valid = count_edges(path, graph)
56     return float(target - valid)
57
58 def fitness_max(path, graph):
59     """Maximisation: Returns number of valid edges. Target = N-1."""
60     return float(count_edges(path, graph))
61
62 # =====
63 # 1b. STATISTICAL ANALYSIS FUNCTIONS
64 # =====
65
66 def cohens_d(group1, group2):
67     """Calculate Cohen's d effect size between two groups."""
68     n1, n2 = len(group1), len(group2)
69     var1, var2 = np.var(group1, ddof=1), np.var(group2, ddof=1)
70     pooled_std = np.sqrt(((n1 - 1) * var1 + (n2 - 1) * var2) / (n1 + n2 - 2))
71     if pooled_std == 0:
72         return 0.0
73     return (np.mean(group1) - np.mean(group2)) / pooled_std
74
75 def interpret_effect_size(d):
76     """Interpret Cohen's d value."""
77     d = abs(d)
78     if d < 0.2:
79         return "negligible"
80     elif d < 0.5:
81         return "small"
82     elif d < 0.8:
83         return "medium"
84     else:
85         return "large"
86
87 def compute_statistical_tests(scores_dict, alpha=0.05, output_file="statistical_results.txt"):
88     """
89     Perform formal statistical tests on algorithm comparison results.
90
91     Args:
92         scores_dict: Dictionary mapping algorithm names to lists of
93             scores
94             alpha: Significance level (default 0.05)
95             output_file: Path to save results (default: statistical_results.
96             txt)
97
98     Returns:
99         Dictionary with test results including p-values and effect sizes
100    """
101    results = {
102        'kruskal_wallis': None,
103        'pairwise_tests': []},

```

```

102     'effect_sizes': [],
103     'confidence_intervals': {}
104 }
105
106 lines = [] # Collect output for file writing
107
108 # 1. Kruskal-Wallis H-test (non-parametric ANOVA)
109 groups = list(scores_dict.values())
110 if len(groups) >= 2:
111     stat, p_value = stats.kruskal(*groups)
112     results['kruskal_wallis'] = {
113         'statistic': stat,
114         'p_value': p_value,
115         'significant': p_value < alpha
116     }
117     lines.append(f"\n--- STATISTICAL ANALYSIS ( ={alpha}) ---")
118     lines.append(f"Kruskal-Wallis H-test: H={stat:.4f}, p={p_value:.4e}")
119     lines.append(f"{'Significant' if p_value < alpha else 'Not significant'} difference between groups")
120
121 # 2. Pairwise Wilcoxon signed-rank tests
122 algo_names = list(scores_dict.keys())
123 lines.append("\nPairwise Wilcoxon signed-rank tests:")
124 for (name1, name2) in combinations(algo_names, 2):
125     scores1, scores2 = scores_dict[name1], scores_dict[name2]
126     try:
127         stat, p_value = stats.wilcoxon(scores1, scores2)
128         significant = p_value < alpha
129     except ValueError:
130         # All differences are zero
131         stat, p_value, significant = 0, 1.0, False
132
133     results['pairwise_tests'].append({
134         'pair': (name1, name2),
135         'statistic': stat,
136         'p_value': p_value,
137         'significant': significant
138     })
139     lines.append(f" {name1} vs {name2}: W={stat:.2f}, p={p_value:.4e} {'*' if significant else ''}")
140
141 # 3. Effect sizes (Cohen's d)
142 lines.append("\nEffect sizes (Cohen's d):")
143 for (name1, name2) in combinations(algo_names, 2):
144     d = cohens_d(scores_dict[name1], scores_dict[name2])
145     interpretation = interpret_effect_size(d)
146     results['effect_sizes'].append({
147         'pair': (name1, name2),
148         'cohens_d': d,
149         'interpretation': interpretation
150     })
151     lines.append(f" {name1} vs {name2}: d={d:.3f} ({interpretation})")
152
153 # 4. Confidence intervals for each algorithm
154 lines.append("\n95% Confidence Intervals:")
155 for name, scores in scores_dict.items():

```

```

156     mean = np.mean(scores)
157     sem = stats.sem(scores)
158     ci = stats.t.interval(0.95, len(scores)-1, loc=mean, scale=sem)
159     results['confidence_intervals'][name] = {'mean': mean, 'ci_low': ci[0], 'ci_high': ci[1]}
160     lines.append(f"  {name}: {mean:.2f} [{ci[0]:.2f}, {ci[1]:.2f}]")
161
162 # Print to console
163 for line in lines:
164     print(line)
165
166 # Write to file
167 with open(output_file, 'a') as f:
168     f.write(f"\n{'='*60}\n")
169     f.write(f"Statistical Analysis Run\n")
170     f.write(f"{'='*60}\n")
171     for line in lines:
172         f.write(line + '\n')
173 print(f"\nResults appended to '{output_file}'")
174
175 return results
176
177 def plot_statistical_comparison(scores_dict, output_prefix="statistical_comparison", title="Algorithm Comparison"):
178     """
179     Create publication-ready statistical comparison visualization.
180
181     Generates:
182     - Box plot with p-value annotations (if statannotations available)
183     - Bar chart with 95% CI error bars
184     - Effect size heatmap
185
186     Args:
187         scores_dict: Dictionary mapping algorithm names to lists of
188             scores
189         output_prefix: Prefix for output files
190         title: Plot title
191     """
192     algo_names = list(scores_dict.keys())
193
194     # Prepare data for plotting
195     plot_data = []
196     for name, scores in scores_dict.items():
197         for score in scores:
198             plot_data.append({'Algorithm': name, 'Broken Edges': score})
199     df = pd.DataFrame(plot_data)
200
201     # Create figure with multiple subplots
202     fig, axes = plt.subplots(1, 3, figsize=(16, 5))
203
204     # --- Panel 1: Box plot with significance annotations ---
205     ax1 = axes[0]
206     sns.boxplot(data=df, x='Algorithm', y='Broken Edges', hue='Algorithm',
207                 legend=False, ax=ax1, palette="viridis")
208     ax1.set_title(f'{title}\n(Lower is Better)')
209     ax1.set_ylim(bottom=-0.5)

```

```

210 # Add p-value annotations if statannotations is available
211 if HAS_STATANNOTATIONS and len(algo_names) >= 2:
212     pairs = list(combinations(algo_names, 2))
213     annotator = Annotator(ax1, pairs, data=df, x='Algorithm', y='Broken Edges')
214     annotator.configure(test='Mann-Whitney', text_format='star', loc='inside',
215                         comparisons_correction=None)
216     annotator.apply_and_annotate()
217
218 # --- Panel 2: Mean with 95% CI error bars ---
219 ax2 = axes[1]
220 means = []
221 ci_lows = []
222 ci_highs = []
223 for name in algo_names:
224     scores = scores_dict[name]
225     mean = np.mean(scores)
226     sem = stats.sem(scores)
227     ci = stats.t.interval(0.95, len(scores)-1, loc=mean, scale=sem)
228     means.append(mean)
229     ci_lows.append(mean - ci[0])
230     ci_highs.append(ci[1] - mean)
231
232 colors = sns.color_palette("viridis", len(algo_names))
233 bars = ax2.bar(algo_names, means, yerr=[ci_lows, ci_highs], capsize=5,
234                 color=colors, edgecolor='black', linewidth=1.5)
235 ax2.set_ylabel('Mean Broken Edges')
236 ax2.set_title('Mean 95% Confidence Interval')
237 ax2.set_ylim(bottom=0)
238
239 # Add value labels on bars
240 for bar, mean in zip(bars, means):
241     ax2.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.5,
242               f'{mean:.2f}', ha='center', va='bottom', fontsize=10,
243               fontweight='bold')
244
245 # --- Panel 3: Effect size heatmap ---
246 ax3 = axes[2]
247 n = len(algo_names)
248 effect_matrix = np.zeros((n, n))
249 for i, name1 in enumerate(algo_names):
250     for j, name2 in enumerate(algo_names):
251         if i != j:
252             d = cohens_d(scores_dict[name1], scores_dict[name2])
253             effect_matrix[i, j] = d
254
255 # Create annotated heatmap
256 mask = np.eye(n, dtype=bool) # Mask diagonal
257 sns.heatmap(effect_matrix, annot=True, fmt='.2f', cmap='RdYlGn_r',
258             xticklabels=algo_names, yticklabels=algo_names,
259             mask=mask, ax=ax3, center=0, vmin=-2, vmax=2,
260             cbar_kws={'label': "Cohen's d"})
261 ax3.set_title("Effect Size (Cohen's d)\n(Row vs Column)")
262 plt.tight_layout()

```

```

263 # Save figure
264 fig_path = f'{output_prefix}_plot.png'
265 plt.savefig(fig_path, dpi=300, bbox_inches='tight')
266 print(f"Saved '{fig_path}'")
267 plt.close()
268
269 # Save summary to CSV
270 summary_data = {
271     'Algorithm': algo_names,
272     'Mean': means,
273     'CI_Lower': [m - cl for m, cl in zip(means, ci_lows)],
274     'CI_Upper': [m + ch for m, ch in zip(means, ci_highs)],
275     'Std': [np.std(scores_dict[name]) for name in algo_names]
276 }
277 summary_df = pd.DataFrame(summary_data)
278 csv_path = f'{output_prefix}_summary.csv'
279 summary_df.to_csv(csv_path, index=False)
280 print(f"Saved '{csv_path}'")
281
282 # =====
283 # 2. OPERATORS (THE OPTIMISATIONS)
284 # =====
285
286
287 def op_swap(genome):
288     """Standard Swap: Exchanges two random nodes."""
289     n = len(genome)
290     i, j = random.sample(range(n), 2)
291     genome[i], genome[j] = genome[j], genome[i]
292     return genome
293
294 def op_inversion(genome):
295     """
296     2-Opt Inversion: Reverses a random sub-segment.
297     OPTIMIZATION: Preserves adjacency better than swap.
298     """
299     n = len(genome)
300     i, j = sorted(random.sample(range(n), 2))
301     # Reverse the segment between i and j
302     genome[i:j+1] = genome[i:j+1][::-1]
303     return genome
304
305 # =====
306 # 3. METAHEURISTIC ALGORITHMS
307 # =====
308
309 def run_simulated_annealing(graph, max_steps=3000, temp0=100.0, operator="swap"):
310     """Simulated Annealing with selectable operator."""
311     nodes = list(graph.nodes)
312     current = get_initial_solution(nodes)
313     best = current.copy()
314
315     current_cost = fitness_min(current, graph)
316     best_cost = current_cost
317     trace = [best_cost]
318     temp = temp0
319

```

```

320     mutate_func = op_inversion if operator == "inversion" else op_swap
321
322     for step in range(max_steps):
323         if best_cost == 0: break
324
325         neighbor = current.copy()
326         neighbor = mutate_func(neighbor)
327
328         neighbor_cost = fitness_min(neighbor, graph)
329         delta = neighbor_cost - current_cost
330
331         if delta < 0 or random.random() < exp(-delta / temp):
332             current = neighbor
333             current_cost = neighbor_cost
334             if current_cost < best_cost:
335                 best = current.copy()
336                 best_cost = current_cost
337
338             trace.append(best_cost)
339             temp *= 0.985
340
341     return best, best_cost, trace
342
343 def run_tabu_search(graph, max_steps=1000, tenure=20, operator="swap"):
344     """Tabu Search with selectable operator."""
345     nodes = list(graph.nodes)
346     current = get_initial_solution(nodes)
347     best = current.copy()
348     best_cost = fitness_min(best, graph)
349
350     tabu_list = []
351     trace = [best_cost]
352     mutate_func = op_inversion if operator == "inversion" else op_swap
353
354     for step in range(max_steps):
355         if best_cost == 0: break
356
357         candidates = []
358         for _ in range(50):
359             cand = current.copy()
360             cand = mutate_func(cand)
361             candidates.append(cand)
362
363         candidates.sort(key=lambda p: fitness_min(p, graph))
364
365         found_move = False
366         for cand in candidates:
367             cand_cost = fitness_min(cand, graph)
368             if cand_cost < best_cost:
369                 current = cand
370                 best = cand
371                 best_cost = cand_cost
372                 found_move = True
373                 break
374             if cand not in tabu_list:
375                 current = cand
376                 found_move = True
377                 break

```

```

378     trace.append(best_cost)
379     if found_move:
380         tabu_list.append(current)
381         if len(tabu_list) > tenure:
382             tabu_list.pop(0)
383
384     return best, best_cost, trace
385
386
387 def run_genetic_algorithm(graph, pop_size=200, generations=1000,
388 cross_rate=0.8, mut_rate=0.3, operator="swap", adaptive=False):
388 """
389     Genetic Algorithm with selectable operator and optional adaptive
390     mutation.
391
392     When adaptive=True:
393         - Calculates population diversity (unique fitness values / pop_size)
394         - Low diversity (< 0.3): INCREASE mutation to escape local optima
395         - High diversity (> 0.7): DECREASE mutation to exploit good
396         solutions
397         - Formula: current_mut = base_rate * (1 + (0.5 - diversity))
398         - Clamped to [0.1, 0.6] range
399     """
400
401     nodes = list(graph.nodes)
402     target = len(nodes) - 1
403
404     mutate_func = op_inversion if operator == "inversion" else op_swap
405     base_mut_rate = mut_rate
406
407
408     def ordered_crossover(p1, p2):
409         size = len(p1)
410         start, end = sorted(random.sample(range(size), 2))
411         child = [None] * size
412         child[start:end] = p1[start:end]
413         current_p2_idx = 0
414         for i in range(size):
415             if child[i] is None:
416                 while p2[current_p2_idx] in child:
417                     current_p2_idx += 1
418                 child[i] = p2[current_p2_idx]
419
420         return child
421
422
423     def calculate_diversity(fitnesses):
424         """Calculate diversity as ratio of unique fitness values."""
425         unique = len(set(fitnesses))
426         return unique / len(fitnesses)
427
428     population = [get_initial_solution(nodes) for _ in range(pop_size)]
429     best_sol = None
430     best_broken = float('inf')
431     trace = []
432     diversity_trace = [] # Track diversity for analysis
433
434     for gen in range(generations):
435         fitnesses = [fitness_max(p, graph) for p in population]
436         best_val = max(fitnesses)
437         current_broken = target - best_val
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
999

```

```

433     if current_broken < best_broken:
434         best_broken = current_broken
435         best_sol = population[fitnesses.index(best_val)]
436
437     trace.append(best_broken)
438     if current_broken == 0: break
439
440     # Adaptive mutation rate adjustment
441     if adaptive:
442         diversity = calculate_diversity(fitnesses)
443         diversity_trace.append(diversity)
444         # Increase mutation when diversity is low, decrease when
445         # high
446         current_mut = base_mut_rate * (1 + (0.5 - diversity))
447         current_mut = max(0.1, min(0.6, current_mut)) # Clamp to
448         [0.1, 0.6]
449     else:
450         current_mut = mut_rate
451
452     next_pop = [best_sol.copy()]
453     while len(next_pop) < pop_size:
454         competitors = random.sample(population, 3)
455         winner = max(competitors, key=lambda p: fitness_max(p, graph
456 ))
457         next_pop.append(winner.copy())
458     population = next_pop
459
460     for i in range(1, pop_size - 1, 2):
461         if random.random() < cross_rate:
462             population[i] = ordered_crossover(population[i],
463 population[i+1])
464
465     for i in range(1, pop_size):
466         if random.random() < current_mut:
467             population[i] = mutate_func(population[i])
468
469     return best_sol, best_broken, trace
470
471 # =====
472 # 4. EXPERIMENT RUNNERS
473 # =====
474
475 def run_batch_experiments(num_nodes, prob):
476     """Runs each algorithm 30 times and plots results."""
477     # Baseline comparison using Standard Swap
478     OPERATOR = "swap"
479
480     g_exp = nx.erdos_renyi_graph(n=num_nodes, p=prob, seed=SEED)
481
482     print(f"\n--- 1. BATCH EXPERIMENT (Baseline 'Swap', N={num_nodes}, p
483 ={prob}, 30 Runs) ---")
484
485     def run_batch(algo_func, name):
486         scores = []
487         times = []
488         successes = 0
489         for _ in range(30):
490             start = time.time()

```

```

486     _, cost, _ = algo_func(g_exp)
487     dur = time.time() - start
488     scores.append(cost)
489     times.append(dur)
490     if cost == 0: successes += 1
491
492     rate = (successes/30)*100
493     print(f"\{name\}: Mean Cost={np.mean(scores):.2f}, Success={rate:.1f}%, Mean Time={np.mean(times):.4f}s")
494     return scores, times
495
496 sa_scores, sa_times = run_batch(lambda g: run_simulated_annealing(g,
497                                 max_steps=3000, operator=OPERATOR), "SA")
498 tabu_scores, tabu_times = run_batch(lambda g: run_tabu_search(g,
499                                 max_steps=1500, operator=OPERATOR), "Tabu")
500 ga_scores, ga_times = run_batch(lambda g: run_genetic_algorithm(g,
501                                 generations=1500, operator=OPERATOR), "GA")
502
503 data = {
504     'Algorithm': ['SA']*30 + ['Tabu']*30 + ['GA']*30,
505     'Broken Edges': sa_scores + tabu_scores + ga_scores,
506     'Time (s)': sa_times + tabu_times + ga_times
507 }
508 df = pd.DataFrame(data)
509
510 fig, axes = plt.subplots(1, 2, figsize=(14, 6))
511
512 sns.boxplot(data=df, x='Algorithm', y='Broken Edges', hue='Algorithm',
513               legend=False, ax=axes[0], palette="viridis")
514 axes[0].set_title(f'Solution Quality (N={num_nodes}, Lower is Better')
515 ) )
516 axes[0].set_ylim(bottom=-0.5)
517
518 sns.boxplot(data=df, x='Algorithm', y='Time (s)', hue='Algorithm',
519               legend=False, ax=axes[1], palette="magma")
520 axes[1].set_title(f'Execution Time (N={num_nodes})')
521 axes[1].set_ylim(bottom=0)
522
523 plt.tight_layout()
524 plt.savefig('experiment_1_batch_results.png', dpi=300, bbox_inches='tight')
525 print("Saved 'experiment_1_batch_results.png'")
526 # plt.show() # Uncomment if running interactively
527
528 def run_optimization_comparison(num_nodes):
529     """Compares Standard 'Swap' vs Optimized 'Inversion' (2-Opt)."""
530     print("\n--- 2. OPTIMISATION EVALUATION (N={num_nodes}) ---")
531     g_opt = nx.erdos_renyi_graph(n=num_nodes, p=0.1, seed=SEED)
532
533     print("Running SA with Standard Swap...")
534     _, _, trace_swap = run_simulated_annealing(g_opt, max_steps=4000,
535                                                 operator="swap")
536
537     print("Running SA with Optimized Inversion (2-Opt)...")
538     _, _, trace_inv = run_simulated_annealing(g_opt, max_steps=4000,
539                                                 operator="inversion")

```

```

535 df_opt = pd.DataFrame({
536     'Step': list(range(len(trace_swap))) + list(range(len(trace_inv)))
537     )),
538     'Broken Edges': trace_swap + trace_inv,
539     'Variant': ['Standard (Swap)'] * len(trace_swap) + ['Optimized
540 (2-Opt)'] * len(trace_inv)
541 })
542
543 plt.figure(figsize=(10, 6))
544 sns.lineplot(data=df_opt, x='Step', y='Broken Edges', hue='Variant',
545 linewidth=2.5)
546 plt.title(f'Impact of Neighborhood Structure: Swap vs 2-Opt (N={num_nodes})')
547 plt.ylabel('Broken Edges (Cost)')
548 plt.xlabel('Iteration Step')
549 plt.ylim(bottom=-0.5)
550
551 plt.savefig('experiment_2_optimization_comparison.png', dpi=300,
552 bbox_inches='tight')
553 print("Saved 'experiment_2_optimization_comparison.png'")
554
555 def run_phase_transition(num_nodes):
556 """
557 Analyzes difficulty vs graph density across ALL algorithms.
558 This provides a comprehensive view of algorithmic limits.
559 """
560 print(f"\n--- 3. PHASE TRANSITION ANALYSIS (N={num_nodes}) ---")
561 # Densities to test
562 densities = [0.05, 0.08, 0.1, 0.12, 0.15, 0.2, 0.25, 0.3]
563
564 # Calculate Theoretical Threshold for THIS N
565 # Koml s & Szemer di theorem: p = (ln(n) + ln(ln(n))) / n
566 if num_nodes > 1:
567     math_threshold = (log(num_nodes) + log(log(num_nodes))) /
568 num_nodes
569     print(f"Theoretical Critical Threshold for N={num_nodes}: p ~ {math_threshold:.3f}")
570 else:
571     math_threshold = 0
572
573 # Store results for plotting
574 results = {'Density': [], 'Success Rate': [], 'Algorithm': []}
575
576 # Test all algorithms to see which one handles the transition best
577 algorithms = {
578     'SA': lambda g: run_simulated_annealing(g, max_steps=2500,
579 operator="inversion"),
580     'Tabu': lambda g: run_tabu_search(g, max_steps=1500, operator="inversion"),
581     'GA': lambda g: run_genetic_algorithm(g, generations=1000,
582 operator="inversion")
583 }
584
585 for p in densities:
586     # Create a graph for this density
587     # Note: Ideally we average over multiple graphs, but for speed
588     # we use one seed per density

```

```

581     # or we generate a fresh one each run. Let's stick to one graph
582     instance per density
583     # but 10 runs per algorithm on it.
584     g = nx.erdos_renyi_graph(n=num_nodes, p=p, seed=999)
585
586     print(f"Testing Density p={p}...")
587
588     for name, algo_func in algorithms.items():
589         successes = 0
590         for _ in range(10): # 10 runs per density/algorithm pair
591             _, cost, _ = algo_func(g)
592             if cost == 0: successes += 1
593
594         rate = (successes/10)*100
595         results['Density'].append(p)
596         results['Success Rate'].append(rate)
597         results['Algorithm'].append(name)
598
599     df_phase = pd.DataFrame(results)
600
601     plt.figure(figsize=(10, 6))
602     sns.lineplot(data=df_phase, x='Density', y='Success Rate', hue='Algorithm', marker='o', linewidth=2)
603
604     # Plot the Theoretical Line
605     plt.axvline(x=math_threshold, color='r', linestyle='--', label=f'Theoretical Limit (p={math_threshold:.2f})')
606
607     plt.title(f'Phase Transition: Experiment vs Theory (N={num_nodes})')
608     plt.xlabel('Graph Density (p)')
609     plt.ylabel('Success Rate (%)')
610     plt.ylim(-5, 105)
611     plt.legend()
612     plt.grid(True)
613
614     plt.savefig('experiment_3_phase_transition.png', dpi=300,
615     bbox_inches='tight')
616     print("Saved 'experiment_3_phase_transition.png'")
617
618 def save_best_graph_html(num_nodes, prob):
619     """Runs GA once and saves the result to HTML."""
620     print("\n--- 4. GENERATING VISUALISATION ---")
621     try:
622         from pyvis.network import Network
623     except ImportError:
624         print("Pyvis not installed. Skipping.")
625         return
626
627     # Use arguments, but if probability is too low, we might not find a
628     # path
629     g = nx.erdos_renyi_graph(n=num_nodes, p=prob, seed=SEED)
630     path, cost, _ = run_genetic_algorithm(g, generations=2000, operator="inversion")
631
632     print(f"Final Path Cost: {cost}")
633
634     net = Network(height="600px", width="100%", cdn_resources='remote')

```

```

633     for n in g.nodes:
634         net.add_node(int(n), label=str(n), color='#97c2fc')
635
636     for u, v in g.edges:
637         net.add_edge(int(u), int(v), color='#e0e0e0', width=1)
638
639     for i in range(len(path) - 1):
640         u, v = int(path[i]), int(path[i+1])
641         if g.has_edge(u, v):
642             net.add_edge(u, v, color='red', width=4)
643         else:
644             net.add_edge(u, v, color='red', width=4, dashes=True)
645
646     net.show("hamiltonian_path.html", notebook=False)
647     print("Saved to 'hamiltonian_path.html'")
648
649 def run_adaptive_comparison(num_nodes, prob=0.1, runs=20):
650     """
651     Compare standard GA vs adaptive mutation GA.
652
653     Generates: experiment_5_adaptive_comparison.png
654     """
655     print(f"\n--- 6. ADAPTIVE MUTATION COMPARISON (N={num_nodes}, p={prob}) ---")
656
657     g = nx.erdos_renyi_graph(n=num_nodes, p=prob, seed=SEED)
658
659     # Run standard GA
660     print("Running Standard GA (fixed mutation)...")
661     standard_scores = []
662     standard_traces = []
663     for _ in range(runs):
664         _, cost, trace = run_genetic_algorithm(g, generations=1000,
665 operator="inversion", adaptive=False)
666         standard_scores.append(cost)
667         standard_traces.append(trace)
668
669     # Run adaptive GA
670     print("Running Adaptive GA (diversity-based mutation)...")
671     adaptive_scores = []
672     adaptive_traces = []
673     for _ in range(runs):
674         _, cost, trace = run_genetic_algorithm(g, generations=1000,
675 operator="inversion", adaptive=True)
676         adaptive_scores.append(cost)
677         adaptive_traces.append(trace)
678
679     # Print summary
680     print(f"\nStandard GA: Mean={np.mean(standard_scores):.2f}, Success={((standard_scores.count(0)/runs)*100:.0f}%")
681     print(f"Adaptive GA: Mean={np.mean(adaptive_scores):.2f}, Success={((adaptive_scores.count(0)/runs)*100:.0f}%")
682
683     # Statistical test and visualization
684     scores_dict = {'Standard GA': standard_scores, 'Adaptive GA':
685     adaptive_scores}
686     compute_statistical_tests(scores_dict, output_file="adaptive_comparison_stats.txt")

```

```

684     plot_statistical_comparison(scores_dict,
685                                 output_prefix="adaptive_comparison_stats"
686                                 ,
687                                 title=f"Standard vs Adaptive GA (N={num_nodes})")
688
689     # Plot convergence comparison (using median trace)
690     def get_median_trace(traces):
691         max_len = max(len(t) for t in traces)
692         padded = [t + [t[-1]] * (max_len - len(t)) for t in traces]
693         return np.median(padded, axis=0)
694
695     median_standard = get_median_trace(standard_traces)
696     median_adaptive = get_median_trace(adaptive_traces)
697
698     plt.figure(figsize=(10, 6))
699     plt.plot(median_standard, label='Standard GA (Fixed Mutation)', linewidth=2)
700     plt.plot(median_adaptive, label='Adaptive GA (Diversity-Based)', linewidth=2, linestyle='--')
701     plt.title(f'GA Convergence: Standard vs Adaptive Mutation (N={num_nodes})')
702     plt.xlabel('Generation')
703     plt.ylabel('Broken Edges (Cost)')
704     plt.ylim(bottom=-0.5)
705     plt.legend()
706     plt.grid(True)
707
708     plt.savefig('experiment_5_adaptive_comparison.png', dpi=300,
709                 bbox_inches='tight')
710     print("Saved 'experiment_5_adaptive_comparison.png'")
711
712     # Save detailed results to file
713     with open('statistical_results.txt', 'a') as f:
714         f.write(f"\n{'*60}\n")
715         f.write(f"Adaptive Mutation Comparison (N={num_nodes}, p={prob})\n")
716         f.write(f"\n{'*60}\n")
717         f.write(f"Standard GA: Mean={np.mean(standard_scores):.2f}, Success={standard_scores.count(0)/runs}*100:.0f%\n")
718         f.write(f"Adaptive GA: Mean={np.mean(adaptive_scores):.2f}, Success={(adaptive_scores.count(0)/runs)*100:.0f}%\n")
719     print("Results appended to 'statistical_results.txt'")
720
721     # =====
722     # MAIN EXECUTION
723     # =====
724
725     if __name__ == "__main__":
726         parser = argparse.ArgumentParser(description="Hamiltonian Path Metaheuristic Analysis")
727         parser.add_argument("-N", "--nodes", type=int, default=50, help="Number of nodes in the graph (default: 50)")
728         parser.add_argument("-p", "--prob", type=float, default=0.1, help="Edge creation probability (default: 0.1)")
729         parser.add_argument("--mode", type=str, default="all",
730                             choices=["all", "batch", "opt", "phase", "visual",
731                                     "stats", "adaptive"],


```

```

729             help="Experiment mode to run individually")
730
731
732     args = parser.parse_args()
733
734     print(f"Running experiments with N={args.nodes} and p={args.prob},
735           Mode={args.mode}")
736
737     # 1. Main Baseline (using 'Swap')
738     if args.mode in ["all", "batch"]:
739         run_batch_experiments(args.nodes, args.prob)
740
741     # 2. The Report Recommendation (Proving 2-Opt is better)
742     if args.mode in ["all", "opt"]:
743         run_optimization_comparison(args.nodes)
744
745     # 3. Physics/Difficulty Analysis (Runs ALL algorithms across
746     # densities)
747     if args.mode in ["all", "phase"]:
748         run_phase_transition(args.nodes)
749
750     # 4. Visual Output
751     if args.mode in ["all", "visual"]:
752         save_best_graph_html(args.nodes, args.prob)
753
754     # 5. Statistical Analysis Mode (batch + stats)
755     if args.mode == "stats":
756         print("\n--- RUNNING BATCH WITH STATISTICAL ANALYSIS ---")
757         OPERATOR = "inversion"
758         g_exp = nx.erdos_renyi_graph(n=args.nodes, p=args.prob, seed=
759                                       SEED)
760
761         def run_batch_for_stats(algo_func, name):
762             scores = []
763             for _ in range(30):
764                 _, cost, _ = algo_func(g_exp)
765                 scores.append(cost)
766             return scores
767
768             sa_scores = run_batch_for_stats(lambda g:
769                 run_simulated_annealing(g, max_steps=3000, operator=OPERATOR), "SA")
770             tabu_scores = run_batch_for_stats(lambda g: run_tabu_search(g,
771                 max_steps=1500, operator=OPERATOR), "Tabu")
772             ga_scores = run_batch_for_stats(lambda g: run_genetic_algorithm(
773                 g, generations=1500, operator=OPERATOR), "GA")
774
775             scores_dict = {'SA': sa_scores, 'Tabu': tabu_scores, 'GA':
776                           ga_scores}
777             compute_statistical_tests(scores_dict)
778             plot_statistical_comparison(scores_dict,
779                                         output_prefix=f"stats_N{args.nodes}
780                                         _p{args.prob}",
781                                         title=f"Algorithm Comparison (N={
782                                         args.nodes}, p={args.prob})")
783
784     # 6. Adaptive Mutation Comparison
785     if args.mode == "adaptive":

```

```
777     run_adaptive_comparison(args.nodes, prob=args.prob)
```

Listing 7: Python 3.14.2