

Metaheuristic approach to the Hamiltonian Path

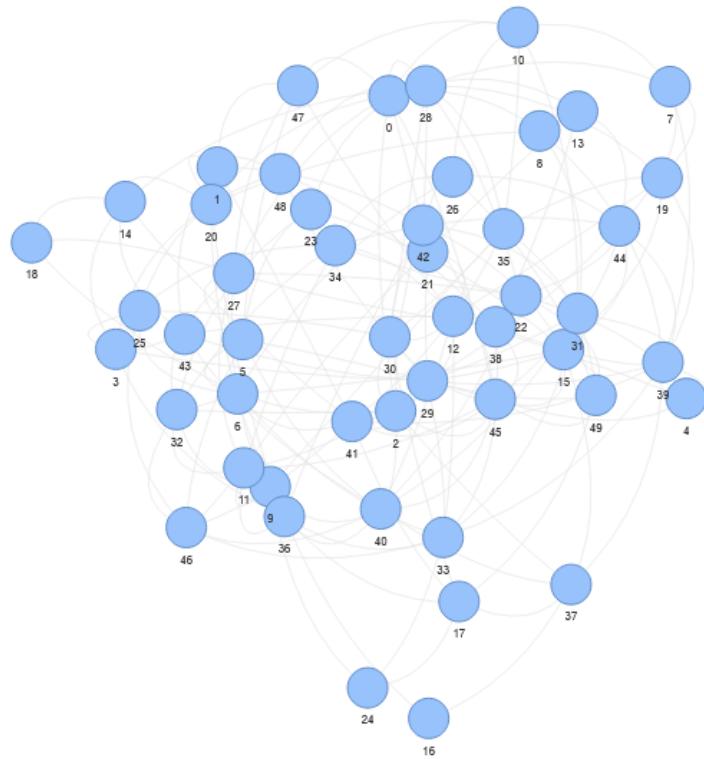


Figure 1: Visualised with Pyvis, depicts the best Hamiltonian path that could be solved for. The GA found using the 2-OPT (inversion) operator in the Erdős-Rényi random graph. Broken edges are depicted in red.

Module: Algorithms and Combinatorial Thinking 2025-2026

Written by:
Christos Christophoros Mitsakopoulos

Supervised by: Dr. Franck Delaplace

January 23, 2026

Contents

1	Formal Definition and Introduction	2
2	Random Generation of the Problem: Erdős-Rényi Random Graph	2
3	Test Environment: Experimental Choices and Reasoning	3
4	Metaheuristic Algorithms: Mode of Action and Code Snippets	4
4.1	Objective Function	4
4.2	Simulated Annealing (SA)	4
4.3	Tabu Search (TS)	5
4.4	Genetic Algorithm (GA)	5
4.4.1	Adaptive Mutation Rate	6
5	Experiment: Comparing the Metaheuristic Algorithms	6
5.1	Batch Testing	6
5.2	Phase Transition	7
6	GA: Comparing Swap Operator to 2-Opt	8
7	Statistical Validation	9
8	Annex	10

1 Formal Definition and Introduction

This report evaluates three metaheuristic approaches: *Simulated Annealing*, *Tabu Search*, and a *Genetic Algorithm*.

To start with, the Hamiltonian Path problem asks if a graph $G = (V, E)$ contains a path that visits every vertex / node $v \in V$ exactly once.

Given that the Hamiltonian Path problem is NP-Complete, as the number of vertices of the vertex set $n = |V|$ increases, the computational complexity required to solve the problem via brute force grows factorially ($O(n!)$). Consequently, exact algorithms become computationally intractable for large n , necessitating the use of metaheuristics – like the ones tested in this report. For this exact reason no evaluation of a brute force approach has been considered in the experiments of this report, as the computational and time complexity required is too high to practically implement.

In the context of bioinformatics, specifically *de novo* genome assembly, this problem is interesting. Arranging DNA reads in an acyclic graph can be modelled as finding a Hamiltonian path that maximises read overlap to create a contiguous sequence. Since this is a bioinformatics course, a genetic algorithm appropriated from the course material (much like all other algorithms implemented) was tested to match the theme of genome assembly heuristics.

Find all relevant results / as well as reproduce the experiment using the `main.py` at the following address: https://github.com/cmitsakopoulos/Delaplace_coursework. The Python script automates all tests demonstrated in this report and moreover, accepts user arguments through the command line interface; intended for tweaking parameters regarding base graph generation. The `-mode` argument allows running specific experiments:

```
1 # Default: runs all standard experiments
2 python main.py --nodes 50 --prob 0.15
3
4 # Specific modes:
5 python main.py --mode batch      # Batch testing only
6 python main.py --mode stats     # Statistical tests (Kruskal-Wallis,
    Wilcoxon)
7 python main.py --mode adaptive # Compare standard vs adaptive mutation
    GA
```

Listing 1: Available experiment modes. The default mode "all" runs batch, optimisation, phase transition, and visualisation experiments.

2 Random Generation of the Problem: Erdős-Rényi Random Graph

To benchmark the metaheuristic algorithms, an Erdős-Rényi (ER) $G(n, p)$ model was used. In which:

- n : The number of vertices in the graph.
- p : The probability that an edge exists between any two distinct vertices.

A larger parameter p will by effect increase the likelihood of finding a Hamiltonian Path, given that the probability of any two nodes having a connecting edge is larger. This was evident when first trying out the `main.py` Python script, where preliminary tests with a p of ≈ 0.3 and $n = 50$ showed that a Hamiltonian Path was indeed mathematically possible to find. It was so much so that all algorithms would converge to a zero cost solution (no broken edges).

```
1 # n_default = 50, p_default = 0.1, seed hardcoded to 42
2 g_exp = nx.erdos_renyi_graph(n=num_nodes, p=prob, seed=SEED)
```

Listing 2: Small snippet of code in which NetworkX in Python can generate an ER graph of chosen parameters, identify that the default options are intended for a "challenging" benchmark for the algorithms. Hardcoding the seed was not a deliberate analytical choice, just to ensure reproducibility of the graph itself – the algorithms are stochastic too...

3 Test Environment: Experimental Choices and Reasoning

To ensure significance of the results, each metaheuristic algorithm was executed for $N = 30$ runs. Relying on the Central Limit Theorem, the sampling distribution of the mean approximates a normal distribution as $N \geq 30$, even if the underlying population distribution is non-Gaussian. While increasing N reduces the Standard Error of the Mean (SEM), the precision improves only with the square root of N (i.e., $SEM = \sigma/\sqrt{N}$).

Additionally, in order to observe the impact on phase transition across all metaheuristics tested, the Komlós–Szemerédi theorem/threshold was used to compute the theoretical p -limit of the problem graph, specifically using the following equation:

$$p_{\text{limit}} = \frac{\ln(n) + \ln(\ln(n))}{n} \quad (1)$$

While $p < p_{\text{limit}}$, the probability of a Hamiltonian path existing approaches zero; thus any algorithm claiming a valid solution is likely "hallucinating", being overconfident (due to verification errors), or has encountered a statistical anomaly (highly unlikely in my testing case). Conversely, for $p > p_{\text{limit}}$, a path almost surely exists: consequently, if an algorithm fails to converge to a zero-cost solution in this region (zero broken edges), it shows that the algorithm is underperforming rather than the problem being impossible.

Another important consideration, was to examine the impact of switching from a "Swap" operator to an **"Inversion" (2-OPT) operator** between testing cases. This was also tested with the GA to demonstrate how 2-OPT can be a positive addition to an already accurate algorithm. Additionally, it was tested with SA and TS to demonstrate that the impact of this change is not algorithm-specific. This phenomenon should become evident when looking at the **"Phase Transition"** experiment of this report, where the 2-OPT operator was used instead of the Swap operator.

To move beyond visual comparisons and quantify whether observed performance differences are statistically meaningful, non-parametric statistical tests were employed. The Kruskal-Wallis H-test determines whether at least one algorithm performs significantly

differently from the others, while pairwise Wilcoxon signed-rank tests identify which specific pairs differ. These non-parametric tests are appropriate because solution quality scores do not follow a normal distribution—algorithms either find optimal solutions (cost = 0) or get stuck at various local optima. Additionally, Cohen’s d effect size quantifies the practical magnitude of differences, distinguishing between statistically significant results and practically meaningful ones. An effect size $|d| \geq 0.8$ indicates a large, practically important difference between algorithms.

4 Metaheuristic Algorithms: Mode of Action and Code Snippets

4.1 Objective Function

The aim of these algorithms is to identify a permutation S of vertices that minimises the number of broken edges in the path. For a graph $G = (V, E)$ and a candidate path $S = [v_1, v_2, \dots, v_n]$, the cost function $C(S)$ is defined as:

$$C(S) = (n - 1) - \sum_{i=1}^{n-1} \mathbb{I}((v_i, v_{i+1}) \in E) \quad (2)$$

Where \mathbb{I} is an indicator function that equals 1 if the edge exists and 0 otherwise. A global optimum is reached when $C(S) = 0$ – a Hamiltonian Path.

4.2 Simulated Annealing (SA)

Simulated Annealing explores solutions by accepting both better or even worse solutions, based on a probability that decreases over time (referred to as parameter T). This logic prevents the algorithm from arriving at a final solution before reaching a truly optimal or near optimal solution. The probability P of accepting a new solution S' with cost difference $\Delta C = C(S') - C(S)$, is governed by the Metropolis criterion (see also the Python implementation below):

$$P(\text{accept}) = \begin{cases} 1 & \text{if } \Delta C < 0 \\ e^{-\frac{\Delta C}{T}} & \text{if } \Delta C \geq 0 \end{cases} \quad (3)$$

```

1 # From main.py: Calculate cost difference
2 delta = neighbor_cost - current_cost
3
4 # Accept if better (delta < 0) OR with probability exp(-delta/T)
5 if delta < 0 or random.random() < exp(-delta / temp):
6     current = neighbor
7     current_cost = neighbor_cost

```

Listing 3: *Metropolis Criterion: As the loop progresses, the temperature T decays geometrically ($T_{k+1} = 0.985 \cdot T_k$), gradually turning the search into a simple greedy descent (hill climbing)*

4.3 Tabu Search (TS)

Tabu Search differs from SA by using a deterministic, memory-based approach. TS explores the immediate neighbourhood of a current solution, then moves to the best available neighbouring solution, even if that neighbouring solution is worse than the current solution.

To prevent cycling (revisiting the same solutions again and again), the algorithm maintains a *Tabu List* – a short-term memory that prevents recent solutions for a specific duration, called tenure.

```
1 # From main.py: Moving to the best candidate not in the Tabu list
2 if cand not in tabu_list:
3     current = cand
4     found_move = True
5
6 # Update memory
7 tabu_list.append(current)
8 if len(tabu_list) > tenure:
9     tabu_list.pop(0) # Remove oldest entry
```

Listing 4: Tabu Search Memory Logic

4.4 Genetic Algorithm (GA)

The Genetic Algorithm attempts to mimic natural selection. Unlike SA and TS, which improve a single solution, GA evolves a population of solutions. The operator for permutation is *Ordered Crossover* (OX1), which is important because standard single-point crossover would result in duplicate or missing vertices / nodes.

The code uses OX1 to preserve the ordering of a sub-segment from one "parent" while filling the remaining slots with genes from the second "parent".

```
1 # From main.py: Preserves sub-segment from parent 1
2 child[start:end] = p1[start:end]
3
4 # Fills remaining slots from parent 2, skipping duplicates
5 for i in range(size):
6     if child[i] is None:
7         while p2[current_p2_idx] in child:
8             current_p2_idx += 1
9         child[i] = p2[current_p2_idx]
```

Listing 5: Ordered Crossover (OX1) Implementation

Using the standard `swap` operator (exchanging two indices) disrupts the adjacency of the path significantly. In comparison, the `inversion` (2-Opt) operator reverses a segment of the path. This is mathematically better for path search problems because it preserves the internal adjacency of the reversed segment, only breaking the two edges at the endpoints of the segment. This distinction is implemented via the ‘`op_inversion`’ function in ‘`main.py`’ and is the primary driver for convergence in denser graphs. A graph is produced at the end during testing to demonstrate the differences of both applications.

$$\text{Swap}(S) \rightarrow \text{High disruption of edges} \quad (4)$$

$$\text{Inversion}(S) \rightarrow \text{Minimal disruption (2-Opt)} \quad (5)$$

4.4.1 Adaptive Mutation Rate

Another improvement over that of the mutation operator (Swap vs 2-Opt), which determines how to mutate a solution in order to develop it, is improving on how often to mutate. With a high fixed mutation rate, the algorithm promotes exploration but disrupts good solutions, while low mutation rates enable exploitation but risk converging early; thereby leading to suboptimal solutions. An adaptive approach addresses this by adjusting the mutation probability μ based on population diversity D :

$$D = \frac{|\mathcal{F}_{\text{unique}}|}{P}, \quad \mu_{\text{adaptive}} = \mu_{\text{base}} \cdot (1 + (0.5 - D)) \quad (6)$$

When $D < 0.5$ (population becoming homogeneous), mutation increases to reintroduce variation. When $D > 0.5$ (healthy diversity), mutation decreases to exploit promising solutions. The rate is clamped to $[0.1, 0.6]$. This self-regulating mechanism helps prevent premature convergence without manual tuning.

```

1 def calculate_diversity(fitnesses):
2     return len(set(fitnesses)) / len(fitnesses)
3
4 # Inside GA loop, per generation:
5 if adaptive:
6     diversity = calculate_diversity(fitnesses)
7     current_mut = base_mut_rate * (1 + (0.5 - diversity))
8     current_mut = max(0.1, min(0.6, current_mut))

```

Listing 6: *Adaptive mutation rate logic from main.py. Diversity is computed per generation, and mutation rate is adjusted accordingly.*

For a fixed mutation rate μ_0 , the change in diversity that is expected per generation matches a drift-selection balance. When diversity is low ($D \rightarrow 0$), the population converges early; such that all individuals occupy similar regions of the search space. The adaptive scheme $\mu_{\text{adaptive}} = \mu_0 \cdot (1.5 - D)$ gives a self-correcting feedback loop: when $D < 0.5$, mutation increases to escape local optima; when $D > 0.5$, mutation decreases to enable for more promising solutions. The equilibrium $D^* = 0.5$ represents an optimal exploration-exploitation balance.

5 Experiment: Comparing the Metaheuristic Algorithms

5.1 Batch Testing

The experiment code in `main.py` has been hardcoded to repeat the stochastic run of each metaheuristic algorithm 30 times, with the user chosen parameters upon initialization through the command line. The results depicted in [Figure 2](#) and [Figure 3](#), were computed with an ER problem graph of $n = 50$ and $p = 0.15$.

Briefly, with reference to [Figure 2](#), the lowest solution quality (highest number of broken edges) is demonstrated by SA, then TS and with the GA, being the best performer. Inversely, the highest execution time per batch run, was demonstrated by GA, followed by TS and lastly SA. These results confirm the expected correlation between computational cost and solution quality (indicating the code works as intended). SA is the fastest because it performs a single $O(1)$ evaluation per step, but its single-trajectory stochasticity struggles to escape bad solutions (local optima) in a solution landscape. TS evaluates a

neighbourhood of size $k = 50$ per step ($O(k)$) (focus on local searches), which linearly increases runtime. Lastly, the GA achieves a near-optimal solution albeit, at the highest computational cost. It maintains diversity through a population-based search ($O(P \cdot N)$ per generation), allowing it to traverse the complex fitness landscape more effectively than the trajectory-based methods.

Nevertheless, the results in [Figure 2](#) clearly show that all algorithms – apart from the GA – are greatly underperforming in this testing case. For an ER graph with $n = 50$, the Komlós Szemerédi limit is $p = 0.106$, considering that the chosen – testing – $p = 0.15$, is greater than the theoretical limit, a solution should be mathematically infeasible; albeit, difficult to obtain.

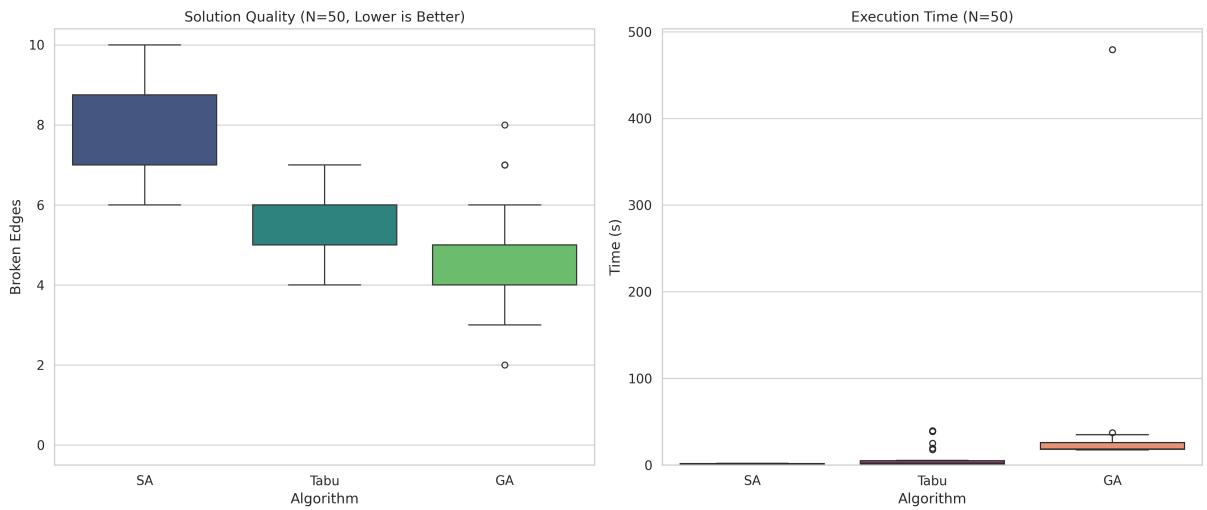


Figure 2: In the box whisker plots depicted, "Solution Quality" (Y-axis: Broken Edges) is shown on the left and "Execution Time" (Y-axis: Time(seconds)) on the right. As discussed prior, due to the low probability of edges between any two nodes, the graph is sparse and the metaheuristic algorithms – while performing differently from one another – are not successful in this testing case. As such, the amount of broken edges seen between algorithms is not unlikely, and the inverse relationship between the number of broken edges and execution time is a clear demonstration of algorithmic complexity – which leads to positive outcomes in problem solving.

5.2 Phase Transition

Briefly, with reference to [Figure 3](#), there appears to be overconfidence of the TS and GA algorithms, which demonstrate a non-zero level of success at a mathematically improbable rate of $p = 0.1$; where $p = 0.1 < p_{\text{limit}} = 0.106$. Given that the difference between p and p_{limit} is $\Delta p = 0.006$, this could be believed to be a statistical anomaly that occurred in this new 30-run test – compared to the previous batch testing scenario. After introducing a new `argparse` argument to the `main.py` file to exclusively re-run the phase transition experiment, the results – over 5 re-runs – led to the same conclusions as in [Figure 3](#). One could argue that due to a **switch from the "Swap" operator to the more optimised "Inversion" (2-Opt) operator**, the GA and TS algorithms are able to find a global optimum at a mathematically improbable rate of $p = 0.1$, as well as be considerably successful at higher graph densities ($p > 0.1$); all the while SA is equally performing better than before (than in [Figure 2](#)).

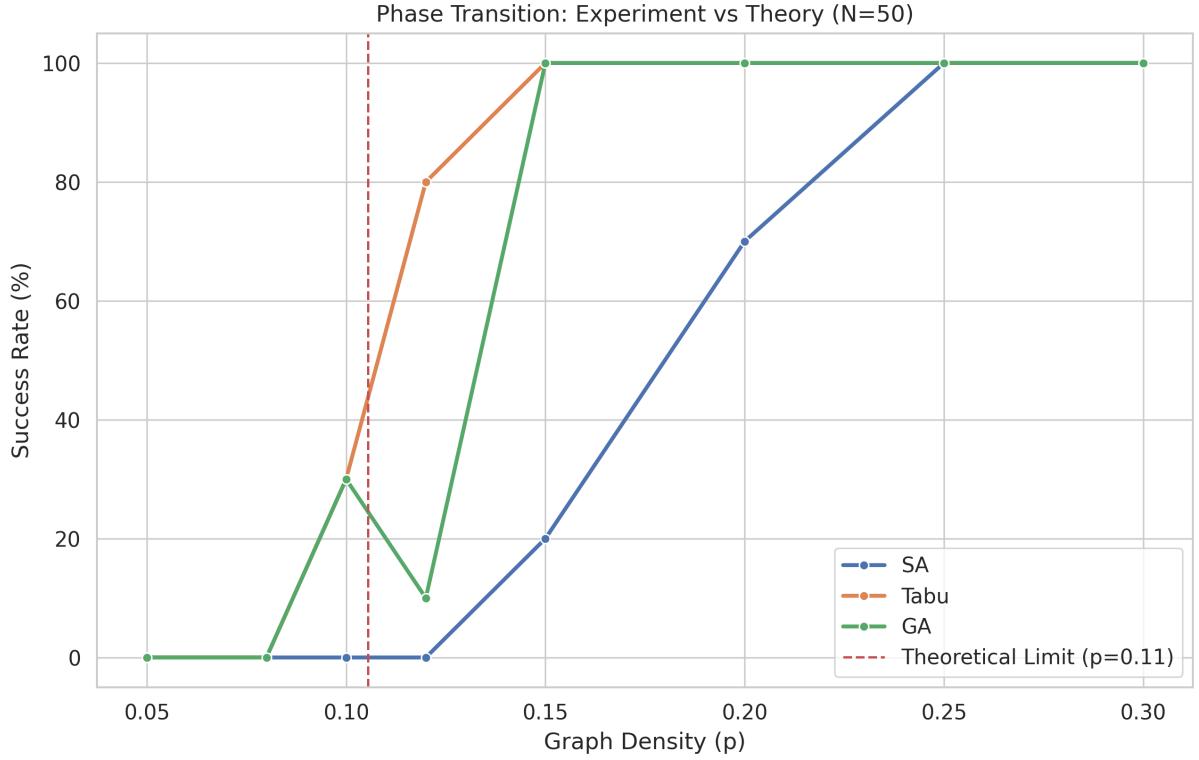


Figure 3: This line chart depicts the phase transition of the proposed problem Erdos-Renyi graph, at a cycling p rate between 0.05 and 0.3. The theoretical limit is computed automatically based on the input problem graph, given the Komlós–Szemerédi formula. Observe overconfidence from the TS and GA algorithms, which demonstrate success at a mathematically improbable rate of p . The success rate of all algorithms is congruent with previous results.

6 GA: Comparing Swap Operator to 2-Opt

Congruent with the phase transition results, it is evident that the switch from the Swap operator (Figure 2) to the Inversion (2-Opt) operator (Figure 3), improves the performance of the GA (see Figure 4). Importantly, looking at the number of broken edges at each iteration, it is clear that the 2-OPT based GA converges at a faster and smoother, or more consistent pace, than the Swap based GA. To briefly reiterate on what was said prior on the report, this phenomenon can be due to two complementary reasons:

Firstly, the 2-OPT (or inversion) operator is less destructive at each iteration. It reverses a segment between indices i and j at each iteration, removing two edges (v_i, v_{i+1}) and (v_j, v_{j+1}) , and replacing them with (v_i, v_j) and (v_{i+1}, v_{j+1}) , whilst preserving the adjacency of all nodes within the segment. As such, the 2-OPT operator enables the algorithm to make small changes which are beneficial at each iteration. Therefore, as can be seen in Figure 4, the GA with the 2-OPT operator improves its global solution at each iteration, whereas the Swap operator will repeatedly plateau/stagnate due to destructive solutions between iterations. Secondly, as mentioned before, the Swap operator exchanges two nodes v_i and v_j arbitrarily, which can affect up to four edges: (v_{i-1}, v_i) , (v_i, v_{i+1}) , (v_{j-1}, v_j) , and (v_j, v_{j+1}) . Instead of making minor improvements locally, the Swap operator will lead the algorithm to more destructive changes at each iteration. In turn, this will cause the search

to stagnate (maintain one solution over multiple iterations) over a certain amount of iterations. While a more beneficial solution can be found accidentally, the algorithm risks getting stuck in local minima; minima which can be avoided by the 2-OPT operator.

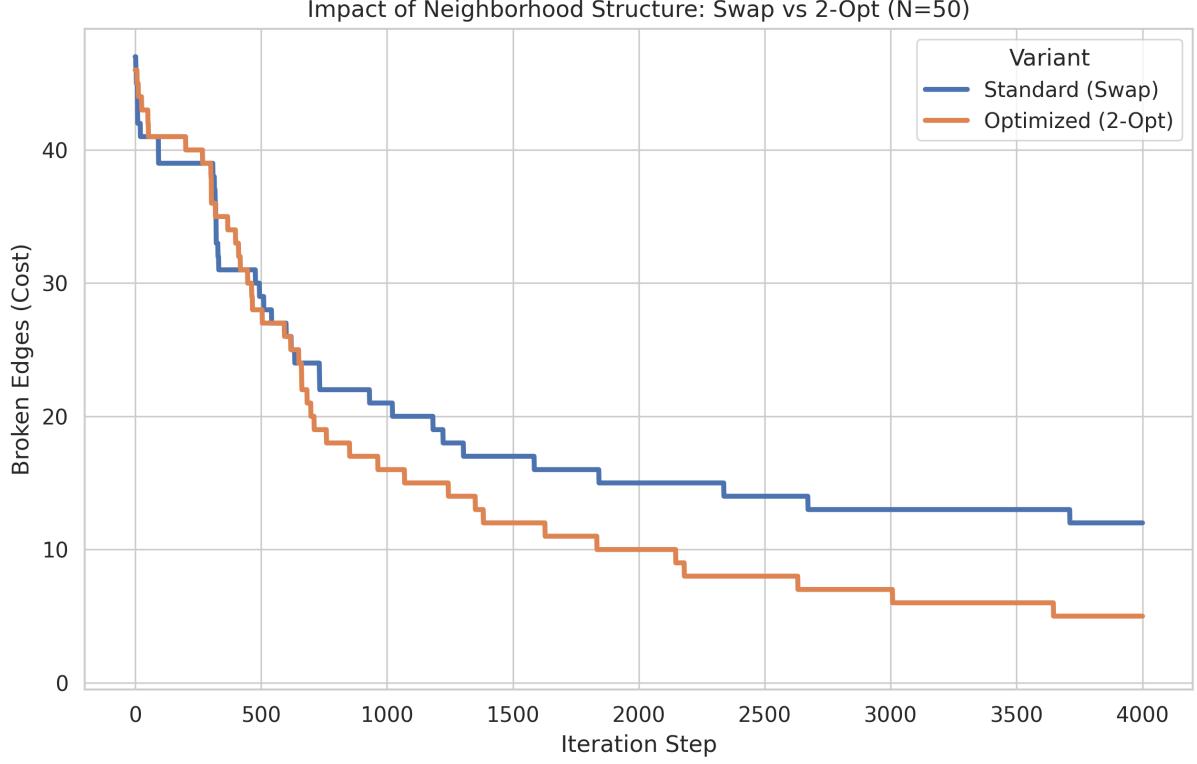


Figure 4: In the depicted line chart, Broken Edges (Y-axis) for which the objective function seeks to minimise (Cost) by maximising the number of edges in the path, is plotted against the number of generations – repeated runs (X-axis). The GA with the Swap operator is shown in blue, while the GA with the Inversion (2-Opt) operator is in orange. Observe the clear disparity between the two operators, in which it's clear that the 2-OPT operator GA converges much faster to a near-global optimum (zero broken edges) than its Swap-GA counterpart.

7 Statistical Validation

To quantify the statistical significance of the performance differences observed in the figures above, the Kruskal-Wallis H-test was applied:

$$H = \frac{12}{N(N+1)} \sum_{i=1}^k \frac{R_i^2}{n_i} - 3(N+1) \quad (7)$$

A significant result ($p < 0.05$) confirms that the algorithms perform differently. Pairwise Wilcoxon signed-rank tests further identify which specific pairs differ:

$$W = \min \left(\sum_{d_i > 0} R_i, \sum_{d_i < 0} R_i \right) \quad (8)$$

Finally, Cohen's d effect size quantifies the magnitude of these differences:

$$d = \frac{\bar{x}_1 - \bar{x}_2}{s_{\text{pooled}}} \quad \text{where} \quad s_{\text{pooled}} = \sqrt{\frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2}} \quad (9)$$

I visualise this statistical analysis through a three-panel figure containing:

- Box plot with Mann-Whitney significance annotations (star notation: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$).
- Mean broken edges with 95% confidence interval error bars.
- Effect size heatmap showing pairwise Cohen's d values.

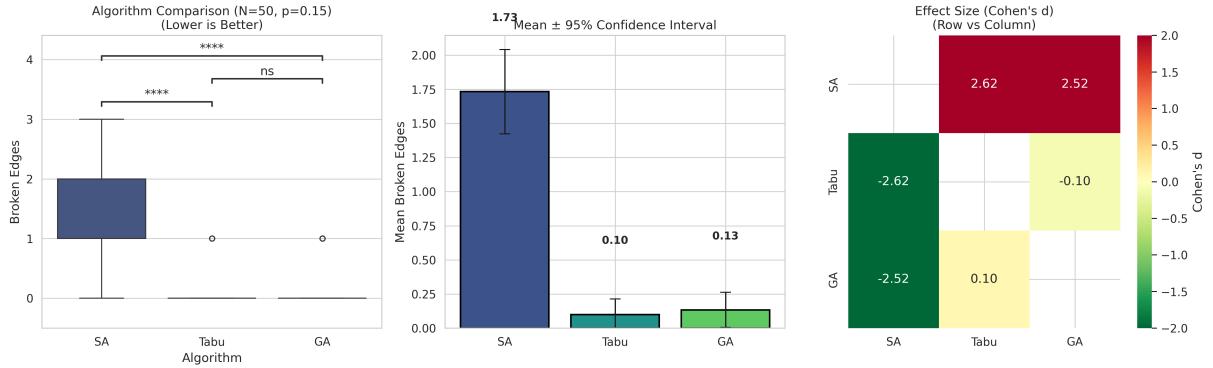


Figure 5: Statistical comparison of SA, Tabu Search, and GA. Left: Box plot with Mann-Whitney significance annotations. Centre: Mean broken edges with 95% CI error bars. Right: Cohen's d effect size heatmap showing pairwise comparisons.

With reference to Figure 5, SA exhibits a mean of $\bar{x}_{SA} = 1.73$ broken edges (95% CI: between 1.42 and 2.04), while Tabu Search ($\bar{x}_{Tabu} = 0.10$, CI: between -0.01 and 0.21) and GA ($\bar{x}_{GA} = 0.13$, CI: between 0.00 and 0.26) achieve near-optimal solutions. Importantly, but expectedly so, the Mann-Whitney tests confirm highly significant differences between SA and both other algorithms ($p < 0.0001$), whereas Tabu vs GA is non-significant. Equally, the effect size heatmap shows large differences between all: $d_{SA \leftrightarrow Tabu} = 2.62$ and $d_{SA \leftrightarrow GA} = 2.52$ (both $|d| > 0.8$), while $d_{Tabu \leftrightarrow GA} = -0.10$ indicates negligible difference. These results statistically confirm that the memory-based (Tabu) and population-based (GA) strategies substantially outperform the single-trajectory stochastic approach (SA) on this graph configuration. However, this strong distinction diminishes at higher graph densities ($p \geq 0.15$), where all algorithms achieve 100% success as shown in the phase transition analysis.

8 Annex

```

1 import random
2 import time
3 import math
4 import itertools
5 import argparse
6 import numpy as np
7 import pandas as pd

```

```

8 import networkx as nx
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11 from math import exp, log
12 from scipy import stats
13 from itertools import combinations
14
15 # Try to import statannotations for publication-ready p-value
16 # annotations
16 try:
17     from statannotations.Annotator import Annotator
18     HAS_STATANNOTATIONS = True
19 except ImportError:
20     HAS_STATANNOTATIONS = False
21     print("Note: Install 'statannotations' for p-value annotations on
22 plots: pip install statannotations")
22
23 # --- CONFIGURATION ---
24 SNS_THEME = "whitegrid"
25 SNS_CONTEXT = "paper"
26 SEED = 42
27
28 # Apply settings
29 random.seed(SEED)
30 np.random.seed(SEED)
31 sns.set_theme(style=SNS_THEME, context=SNS_CONTEXT, font_scale=1.2)
32
33 # =====
34 # 1. CORE HELPER & FITNESS FUNCTIONS
35 # =====
36
37 def get_initial_solution(nodes):
38     """Returns a random permutation of nodes."""
39     p = list(nodes).copy()
40     random.shuffle(p)
41     return p
42
43 def count_edges(path, graph):
44     """Counts valid edges in the path for a specific graph."""
45     edges = 0
46     # Optimization: Iterate via index to avoid creating new lists
47     for i in range(len(path) - 1):
48         if graph.has_edge(path[i], path[i+1]):
49             edges += 1
50     return edges
51
52 def fitness_min(path, graph):
53     """Minimisation: Returns number of broken edges. Target = 0."""
54     target = len(path) - 1
55     valid = count_edges(path, graph)
56     return float(target - valid)
57
58 def fitness_max(path, graph):
59     """Maximisation: Returns number of valid edges. Target = N-1."""
60     return float(count_edges(path, graph))
61
62 # =====
63 # 1b. STATISTICAL ANALYSIS FUNCTIONS

```

```

64 # =====
65
66 def cohens_d(group1, group2):
67     """Calculate Cohen's d effect size between two groups."""
68     n1, n2 = len(group1), len(group2)
69     var1, var2 = np.var(group1, ddof=1), np.var(group2, ddof=1)
70     pooled_std = np.sqrt(((n1 - 1) * var1 + (n2 - 1) * var2) / (n1 + n2 - 2))
71     if pooled_std == 0:
72         return 0.0
73     return (np.mean(group1) - np.mean(group2)) / pooled_std
74
75 def interpret_effect_size(d):
76     """Interpret Cohen's d value."""
77     d = abs(d)
78     if d < 0.2:
79         return "negligible"
80     elif d < 0.5:
81         return "small"
82     elif d < 0.8:
83         return "medium"
84     else:
85         return "large"
86
87 def compute_statistical_tests(scores_dict, alpha=0.05, output_file="statistical_results.txt"):
88     """
89         Perform formal statistical tests on algorithm comparison results.
90
91     Args:
92         scores_dict: Dictionary mapping algorithm names to lists of scores
93         alpha: Significance level (default 0.05)
94         output_file: Path to save results (default: statistical_results.txt)
95
96     Returns:
97         Dictionary with test results including p-values and effect sizes
98     """
99     results = {
100         'kruskal_wallis': None,
101         'pairwise_tests': [],
102         'effect_sizes': [],
103         'confidence_intervals': {}
104     }
105
106     lines = [] # Collect output for file writing
107
108     # 1. Kruskal-Wallis H-test (non-parametric ANOVA)
109     groups = list(scores_dict.values())
110     if len(groups) >= 2:
111         stat, p_value = stats.kruskal(*groups)
112         results['kruskal_wallis'] = {
113             'statistic': stat,
114             'p_value': p_value,
115             'significant': p_value < alpha
116         }
117         lines.append(f"\n--- STATISTICAL ANALYSIS ( ={alpha}) ---")

```

```

118     lines.append(f"Kruskal-Wallis H-test: H={stat:.4f}, p={p_value
119         :.4e}")
120     lines.append(f"      {'Significant' if p_value < alpha else 'Not
121         significant'} difference between groups")
122
123 # 2. Pairwise Wilcoxon signed-rank tests
124 algo_names = list(scores_dict.keys())
125 lines.append("\nPairwise Wilcoxon signed-rank tests:")
126 for (name1, name2) in combinations(algo_names, 2):
127     scores1, scores2 = scores_dict[name1], scores_dict[name2]
128     try:
129         stat, p_value = stats.wilcoxon(scores1, scores2)
130         significant = p_value < alpha
131     except ValueError:
132         # All differences are zero
133         stat, p_value, significant = 0, 1.0, False
134
135     results['pairwise_tests'].append({
136         'pair': (name1, name2),
137         'statistic': stat,
138         'p_value': p_value,
139         'significant': significant
140     })
141     lines.append(f" {name1} vs {name2}: W={stat:.2f}, p={p_value:.4
142         e} {'*' if significant else ''}")
143
144 # 3. Effect sizes (Cohen's d)
145 lines.append("\nEffect sizes (Cohen's d):")
146 for (name1, name2) in combinations(algo_names, 2):
147     d = cohens_d(scores_dict[name1], scores_dict[name2])
148     interpretation = interpret_effect_size(d)
149     results['effect_sizes'].append({
150         'pair': (name1, name2),
151         'cohens_d': d,
152         'interpretation': interpretation
153     })
154     lines.append(f" {name1} vs {name2}: d={d:.3f} ({interpretation
155         })")
156
157 # 4. Confidence intervals for each algorithm
158 lines.append("\n95% Confidence Intervals:")
159 for name, scores in scores_dict.items():
160     mean = np.mean(scores)
161     sem = stats.sem(scores)
162     ci = stats.t.interval(0.95, len(scores)-1, loc=mean, scale=sem)
163     results['confidence_intervals'][name] = {'mean': mean, 'ci_low':
164         ci[0], 'ci_high': ci[1]}
165     lines.append(f" {name}: {mean:.2f} [{ci[0]:.2f}, {ci[1]:.2f}]")
166
167 # Print to console
168 for line in lines:
169     print(line)
170
171 # Write to file
172 with open(output_file, 'a') as f:
173     f.write(f"\n{'='*60}\n")
174     f.write(f"Statistical Analysis Run\n")
175     f.write(f"{'='*60}\n")

```

```

171     for line in lines:
172         f.write(line + '\n')
173     print(f"\nResults appended to '{output_file}'")
174
175     return results
176
177 def plot_statistical_comparison(scores_dict, output_prefix="statistical_comparison", title="Algorithm Comparison"):
178     """
179     Create publication-ready statistical comparison visualization.
180
181     Generates:
182     - Box plot with p-value annotations (if statannotations available)
183     - Bar chart with 95% CI error bars
184     - Effect size heatmap
185
186     Args:
187         scores_dict: Dictionary mapping algorithm names to lists of
188             scores
189         output_prefix: Prefix for output files
190         title: Plot title
191     """
192
193     algo_names = list(scores_dict.keys())
194
195     # Prepare data for plotting
196     plot_data = []
197     for name, scores in scores_dict.items():
198         for score in scores:
199             plot_data.append({'Algorithm': name, 'Broken Edges': score})
200     df = pd.DataFrame(plot_data)
201
202     # Create figure with multiple subplots
203     fig, axes = plt.subplots(1, 3, figsize=(16, 5))
204
205     # --- Panel 1: Box plot with significance annotations ---
206     ax1 = axes[0]
207     sns.boxplot(data=df, x='Algorithm', y='Broken Edges', hue='Algorithm',
208                 legend=False, ax=ax1, palette="viridis")
209     ax1.set_title(f'{title}\n(Lower is Better)')
210     ax1.set_ylim(bottom=-0.5)
211
212     # Add p-value annotations if statannotations is available
213     if HAS_STATANNOTATIONS and len(algo_names) >= 2:
214         pairs = list(combinations(algo_names, 2))
215         annotator = Annotator(ax1, pairs, data=df, x='Algorithm', y='Broken Edges')
216         annotator.configure(test='Mann-Whitney', text_format='star', loc='inside',
217                             comparisons_correction=None)
218         annotator.apply_and_annotate()
219
220     # --- Panel 2: Mean with 95% CI error bars ---
221     ax2 = axes[1]
222     means = []
223     ci_lows = []
224     ci_highs = []
225     for name in algo_names:

```

```

224     scores = scores_dict[name]
225     mean = np.mean(scores)
226     sem = stats.sem(scores)
227     ci = stats.t.interval(0.95, len(scores)-1, loc=mean, scale=sem)
228     means.append(mean)
229     ci_lows.append(mean - ci[0])
230     ci_highs.append(ci[1] - mean)
231
232     colors = sns.color_palette("viridis", len(algo_names))
233     bars = ax2.bar(algo_names, means, yerr=[ci_lows, ci_highs], capsize
234 =5,
235             color=colors, edgecolor='black', linewidth=1.5)
236     ax2.set_ylabel('Mean Broken Edges')
237     ax2.set_title('Mean 95% Confidence Interval')
238     ax2.set_ylim(bottom=0)
239
240     # Add value labels on bars
241     for bar, mean in zip(bars, means):
242         ax2.text(bar.get_x() + bar.get_width()/2, bar.get_height() +
243 0.5,
244             f'{mean:.2f}', ha='center', va='bottom', fontsize=10,
245             fontweight='bold')
246
247     # --- Panel 3: Effect size heatmap ---
248     ax3 = axes[2]
249     n = len(algo_names)
250     effect_matrix = np.zeros((n, n))
251     for i, name1 in enumerate(algo_names):
252         for j, name2 in enumerate(algo_names):
253             if i != j:
254                 d = cohens_d(scores_dict[name1], scores_dict[name2])
255                 effect_matrix[i, j] = d
256
257     # Create annotated heatmap
258     mask = np.eye(n, dtype=bool) # Mask diagonal
259     sns.heatmap(effect_matrix, annot=True, fmt='.2f', cmap='RdYlGn_r',
260                 xticklabels=algo_names, yticklabels=algo_names,
261                 mask=mask, ax=ax3, center=0, vmin=-2, vmax=2,
262                 cbar_kws={'label': "Cohen's d"})
263     ax3.set_title("Effect Size (Cohen's d)\n(Row vs Column)")
264
265     plt.tight_layout()
266
267     # Save figure
268     fig_path = f'{output_prefix}_plot.png'
269     plt.savefig(fig_path, dpi=300, bbox_inches='tight')
270     print(f"Saved '{fig_path}'")
271     plt.close()
272
273     # Save summary to CSV
274     summary_data = {
275         'Algorithm': algo_names,
276         'Mean': means,
277         'CI_Lower': [m - cl for m, cl in zip(means, ci_lows)],
278         'CI_Upper': [m + ch for m, ch in zip(means, ci_highs)],
279         'Std': [np.std(scores_dict[name]) for name in algo_names]
280     }
281     summary_df = pd.DataFrame(summary_data)

```

```

279     csv_path = f'{output_prefix}_summary.csv'
280     summary_df.to_csv(csv_path, index=False)
281     print(f"Saved '{csv_path}'")
282
283 # =====
284 # 2. OPERATORS (THE OPTIMISATIONS)
285 # =====
286
287 def op_swap(genome):
288     """Standard Swap: Exchanges two random nodes."""
289     n = len(genome)
290     i, j = random.sample(range(n), 2)
291     genome[i], genome[j] = genome[j], genome[i]
292     return genome
293
294 def op_inversion(genome):
295     """
296     2-Opt Inversion: Reverses a random sub-segment.
297     OPTIMIZATION: Preserves adjacency better than swap.
298     """
299     n = len(genome)
300     i, j = sorted(random.sample(range(n), 2))
301     # Reverse the segment between i and j
302     genome[i:j+1] = genome[i:j+1][::-1]
303     return genome
304
305 # =====
306 # 3. METAHEURISTIC ALGORITHMS
307 # =====
308
309 def run_simulated_annealing(graph, max_steps=3000, temp0=100.0, operator="swap"):
310     """Simulated Annealing with selectable operator."""
311     nodes = list(graph.nodes)
312     current = get_initial_solution(nodes)
313     best = current.copy()
314
315     current_cost = fitness_min(current, graph)
316     best_cost = current_cost
317     trace = [best_cost]
318     temp = temp0
319
320     mutate_func = op_inversion if operator == "inversion" else op_swap
321
322     for step in range(max_steps):
323         if best_cost == 0: break
324
325         neighbor = current.copy()
326         neighbor = mutate_func(neighbor)
327
328         neighbor_cost = fitness_min(neighbor, graph)
329         delta = neighbor_cost - current_cost
330
331         if delta < 0 or random.random() < exp(-delta / temp):
332             current = neighbor
333             current_cost = neighbor_cost
334             if current_cost < best_cost:
335                 best = current.copy()

```

```

336         best_cost = current_cost
337
338     trace.append(best_cost)
339     temp *= 0.985
340
341     return best, best_cost, trace
342
343 def run_tabu_search(graph, max_steps=1000, tenure=20, operator="swap"):
344     """Tabu Search with selectable operator."""
345     nodes = list(graph.nodes)
346     current = get_initial_solution(nodes)
347     best = current.copy()
348     best_cost = fitness_min(best, graph)
349
350     tabu_list = []
351     trace = [best_cost]
352     mutate_func = op_inversion if operator == "inversion" else op_swap
353
354     for step in range(max_steps):
355         if best_cost == 0: break
356
357         candidates = []
358         for _ in range(50):
359             cand = current.copy()
360             cand = mutate_func(cand)
361             candidates.append(cand)
362
363         candidates.sort(key=lambda p: fitness_min(p, graph))
364
365         found_move = False
366         for cand in candidates:
367             cand_cost = fitness_min(cand, graph)
368             if cand_cost < best_cost:
369                 current = cand
370                 best = cand
371                 best_cost = cand_cost
372                 found_move = True
373                 break
374             if cand not in tabu_list:
375                 current = cand
376                 found_move = True
377                 break
378
379         trace.append(best_cost)
380         if found_move:
381             tabu_list.append(current)
382             if len(tabu_list) > tenure:
383                 tabu_list.pop(0)
384
385     return best, best_cost, trace
386
387 def run_genetic_algorithm(graph, pop_size=200, generations=1000,
388 cross_rate=0.8, mut_rate=0.3, operator="swap", adaptive=False):
389     """
390     Genetic Algorithm with selectable operator and optional adaptive
391     mutation.
392
393     When adaptive=True:

```

```

392     - Calculates population diversity (unique fitness values / pop_size)
393     - Low diversity (< 0.3): INCREASE mutation to escape local optima
394     - High diversity (> 0.7): DECREASE mutation to exploit good
395     solutions
396     - Formula: current_mut = base_rate * (1 + (0.5 - diversity))
397     - Clamped to [0.1, 0.6] range
398     """
399
400     nodes = list(graph.nodes)
401     target = len(nodes) - 1
402
403     mutate_func = op_inversion if operator == "inversion" else op_swap
404     base_mut_rate = mut_rate
405
406     def ordered_crossover(p1, p2):
407         size = len(p1)
408         start, end = sorted(random.sample(range(size), 2))
409         child = [None] * size
410         child[start:end] = p1[start:end]
411         current_p2_idx = 0
412         for i in range(size):
413             if child[i] is None:
414                 while p2[current_p2_idx] in child:
415                     current_p2_idx += 1
416                 child[i] = p2[current_p2_idx]
417         return child
418
419     def calculate_diversity(fitnesses):
420         """Calculate diversity as ratio of unique fitness values."""
421         unique = len(set(fitnesses))
422         return unique / len(fitnesses)
423
424     population = [get_initial_solution(nodes) for _ in range(pop_size)]
425     best_sol = None
426     best_broken = float('inf')
427     trace = []
428     diversity_trace = [] # Track diversity for analysis
429
430     for gen in range(generations):
431         fitnesses = [fitness_max(p, graph) for p in population]
432         best_val = max(fitnesses)
433         current_broken = target - best_val
434
435         if current_broken < best_broken:
436             best_broken = current_broken
437             best_sol = population[fitnesses.index(best_val)]
438
439         trace.append(best_broken)
440         if current_broken == 0: break
441
442         # Adaptive mutation rate adjustment
443         if adaptive:
444             diversity = calculate_diversity(fitnesses)
445             diversity_trace.append(diversity)
446             # Increase mutation when diversity is low, decrease when
447             # high
448             current_mut = base_mut_rate * (1 + (0.5 - diversity))
449             current_mut = max(0.1, min(0.6, current_mut)) # Clamp to
450             [0.1, 0.6]

```

```

447     else:
448         current_mut = mut_rate
449
450     next_pop = [best_sol.copy()]
451     while len(next_pop) < pop_size:
452         competitors = random.sample(population, 3)
453         winner = max(competitors, key=lambda p: fitness_max(p, graph
454 ))
455         next_pop.append(winner.copy())
456     population = next_pop
457
458     for i in range(1, pop_size - 1, 2):
459         if random.random() < cross_rate:
460             population[i] = ordered_crossover(population[i],
461 population[i+1])
462
463     for i in range(1, pop_size):
464         if random.random() < current_mut:
465             population[i] = mutate_func(population[i])
466
467     return best_sol, best_broken, trace
468
469 # =====
470 # 4. EXPERIMENT RUNNERS
471 # =====
472
473 def run_batch_experiments(num_nodes, prob):
474     """Runs each algorithm 30 times and plots results."""
475     # Baseline comparison using Standard Swap
476     OPERATOR = "swap"
477
478     g_exp = nx.erdos_renyi_graph(n=num_nodes, p=prob, seed=SEED)
479
480     print(f"\n--- 1. BATCH EXPERIMENT (Baseline 'Swap', N={num_nodes}, p={prob}, 30 Runs) ---")
481
482     def run_batch(algo_func, name):
483         scores = []
484         times = []
485         successes = 0
486         for _ in range(30):
487             start = time.time()
488             _, cost, _ = algo_func(g_exp)
489             dur = time.time() - start
490             scores.append(cost)
491             times.append(dur)
492             if cost == 0: successes += 1
493
494         rate = (successes/30)*100
495         print(f"{name}: Mean Cost={np.mean(scores):.2f}, Success={rate:.1f}%, Mean Time={np.mean(times):.4f}s")
496         return scores, times
497
498     sa_scores, sa_times = run_batch(lambda g: run_simulated_annealing(g,
499     max_steps=3000, operator=OPERATOR), "SA")
500     tabu_scores, tabu_times = run_batch(lambda g: run_tabu_search(g,
501     max_steps=1500, operator=OPERATOR), "Tabu")
502     ga_scores, ga_times = run_batch(lambda g: run_genetic_algorithm(g,
503

```

```

generations=1500, operator=OPERATOR), "GA")

499
500 data = {
501     'Algorithm': ['SA']*30 + ['Tabu']*30 + ['GA']*30,
502     'Broken Edges': sa_scores + tabu_scores + ga_scores,
503     'Time (s)': sa_times + tabu_times + ga_times
504 }
505 df = pd.DataFrame(data)
506
507 fig, axes = plt.subplots(1, 2, figsize=(14, 6))
508
509 sns.boxplot(data=df, x='Algorithm', y='Broken Edges', hue='Algorithm',
510             legend=False, ax=axes[0], palette="viridis")
511 axes[0].set_title(f'Solution Quality (N={num_nodes}, Lower is Better')
512             )
513 axes[0].set_ylim(bottom=-0.5)
514
515 sns.boxplot(data=df, x='Algorithm', y='Time (s)', hue='Algorithm',
516             legend=False, ax=axes[1], palette="magma")
517 axes[1].set_title(f'Execution Time (N={num_nodes})')
518 axes[1].set_ylim(bottom=0)
519
520 plt.tight_layout()
521 plt.savefig('experiment_1_batch_results.png', dpi=300, bbox_inches='tight')
522 print("Saved 'experiment_1_batch_results.png'")
523 # plt.show() # Uncomment if running interactively
524
524 def run_optimization_comparison(num_nodes):
525     """Compares Standard 'Swap' vs Optimized 'Inversion' (2-Opt)."""
526     print(f"\n--- 2. OPTIMISATION EVALUATION (N={num_nodes}) ---")
527     g_opt = nx.erdos_renyi_graph(n=num_nodes, p=0.1, seed=SEED)
528
529     print("Running SA with Standard Swap...")
530     _, _, trace_swap = run_simulated_annealing(g_opt, max_steps=4000,
531                                                 operator="swap")
532
533     print("Running SA with Optimized Inversion (2-Opt)...")
534     _, _, trace_inv = run_simulated_annealing(g_opt, max_steps=4000,
535                                                 operator="inversion")
536
537     df_opt = pd.DataFrame({
538         'Step': list(range(len(trace_swap))) + list(range(len(trace_inv))),
539         'Broken Edges': trace_swap + trace_inv,
540         'Variant': ['Standard (Swap)'] * len(trace_swap) + ['Optimized
541 (2-Opt)'] * len(trace_inv)
542     })
543
544     plt.figure(figsize=(10, 6))
545     sns.lineplot(data=df_opt, x='Step', y='Broken Edges', hue='Variant',
546                 linewidth=2.5)
547     plt.title(f'Impact of Neighborhood Structure: Swap vs 2-Opt (N={num_nodes})')
548     plt.ylabel('Broken Edges (Cost)')
549     plt.xlabel('Iteration Step')
550     plt.ylim(bottom=-0.5)

```

```

547
548     plt.savefig('experiment_2_optimization_comparison.png', dpi=300,
549     bbox_inches='tight')
550     print("Saved 'experiment_2_optimization_comparison.png'")
551
551 def run_phase_transition(num_nodes):
552     """
553     Analyzes difficulty vs graph density across ALL algorithms.
554     This provides a comprehensive view of algorithmic limits.
555     """
556     print(f"\n--- 3. PHASE TRANSITION ANALYSIS (N={num_nodes}) ---")
557     # Densities to test
558     densities = [0.05, 0.08, 0.1, 0.12, 0.15, 0.2, 0.25, 0.3]
559
560     # Calculate Theoretical Threshold for THIS N
561     # Koml's & Szemer di theorem: p = (ln(n) + ln(ln(n))) / n
562     if num_nodes > 1:
563         math_threshold = (log(num_nodes) + log(log(num_nodes))) / num_nodes
564         print(f"Theoretical Critical Threshold for N={num_nodes}: p ~ {math_threshold:.3f}")
565     else:
566         math_threshold = 0
567
568     # Store results for plotting
569     results = {'Density': [], 'Success Rate': [], 'Algorithm': []}
570
571     # Test all algorithms to see which one handles the transition best
572     algorithms = {
573         'SA': lambda g: run_simulated_annealing(g, max_steps=2500,
574 operator="inversion"),
574         'Tabu': lambda g: run_tabu_search(g, max_steps=1500, operator="inversion"),
575         'GA': lambda g: run_genetic_algorithm(g, generations=1000,
575 operator="inversion")
576     }
577
578     for p in densities:
579         # Create a graph for this density
580         # Note: Ideally we average over multiple graphs, but for speed
581         # we use one seed per density
582         # or we generate a fresh one each run. Let's stick to one graph
583         # instance per density
584         # but 10 runs per algorithm on it.
585         g = nx.erdos_renyi_graph(n=num_nodes, p=p, seed=999)
586
587         print(f"Testing Density p={p}...")
588
589         for name, algo_func in algorithms.items():
590             successes = 0
591             for _ in range(10): # 10 runs per density/algorithm pair
592                 _, cost, _ = algo_func(g)
593                 if cost == 0: successes += 1
594
595             rate = (successes/10)*100
596             results['Density'].append(p)
597             results['Success Rate'].append(rate)
598             results['Algorithm'].append(name)

```

```

597 df_phase = pd.DataFrame(results)
598
599 plt.figure(figsize=(10, 6))
600 sns.lineplot(data=df_phase, x='Density', y='Success Rate', hue='Algorithm', marker='o', linewidth=2)
601
602 # Plot the Theoretical Line
603 plt.axvline(x=math_threshold, color='r', linestyle='--', label=f'Theoretical Limit (p={math_threshold:.2f})')
604
605 plt.title(f'Phase Transition: Experiment vs Theory (N={num_nodes})')
606 plt.xlabel('Graph Density (p)')
607 plt.ylabel('Success Rate (%)')
608 plt.ylim(-5, 105)
609 plt.legend()
610 plt.grid(True)
611
612 plt.savefig('experiment_3_phase_transition.png', dpi=300,
613 bbox_inches='tight')
614 print("Saved 'experiment_3_phase_transition.png'")
615
616 def save_best_graph_html(num_nodes, prob):
617     """Runs GA once and saves the result to HTML."""
618     print("\n--- 4. GENERATING VISUALISATION ---")
619     try:
620         from pyvis.network import Network
621     except ImportError:
622         print("Pyvis not installed. Skipping.")
623         return
624
625     # Use arguments, but if probability is too low, we might not find a
626     # path
627     g = nx.erdos_renyi_graph(n=num_nodes, p=prob, seed=SEED)
628     path, cost, _ = run_genetic_algorithm(g, generations=2000, operator="inversion")
629
630     print(f"Final Path Cost: {cost}")
631
632     net = Network(height="600px", width="100%", cdn_resources='remote')
633
634     for n in g.nodes:
635         net.add_node(int(n), label=str(n), color="#97c2fc")
636
637     for u, v in g.edges:
638         net.add_edge(int(u), int(v), color="#e0e0e0", width=1)
639
640     for i in range(len(path) - 1):
641         u, v = int(path[i]), int(path[i+1])
642         if g.has_edge(u, v):
643             net.add_edge(u, v, color='red', width=4)
644         else:
645             net.add_edge(u, v, color='red', width=4, dashes=True)
646
647     net.show("hamiltonian_path.html", notebook=False)
648     print("Saved to 'hamiltonian_path.html'")
649
650 def run_adaptive_comparison(num_nodes, prob=0.1, runs=20):

```

```

650 """
651     Compare standard GA vs adaptive mutation GA.
652
653 Generates: experiment_5_adaptive_comparison.png
654 """
655 print(f"\n--- 6. ADAPTIVE MUTATION COMPARISON (N={num_nodes}, p={prob}) ---")
656
657 g = nx.erdos_renyi_graph(n=num_nodes, p=prob, seed=SEED)
658
659 # Run standard GA
660 print("Running Standard GA (fixed mutation)...")
661 standard_scores = []
662 standard_traces = []
663 for _ in range(runs):
664     _, cost, trace = run_genetic_algorithm(g, generations=1000,
665 operator="inversion", adaptive=False)
666     standard_scores.append(cost)
667     standard_traces.append(trace)
668
669 # Run adaptive GA
670 print("Running Adaptive GA (diversity-based mutation)...")
671 adaptive_scores = []
672 adaptive_traces = []
673 for _ in range(runs):
674     _, cost, trace = run_genetic_algorithm(g, generations=1000,
675 operator="inversion", adaptive=True)
676     adaptive_scores.append(cost)
677     adaptive_traces.append(trace)
678
679 # Print summary
680 print(f"\nStandard GA: Mean={np.mean(standard_scores):.2f}, Success={((standard_scores.count(0)/runs)*100:.0f}%)")
681 print(f"Adaptive GA: Mean={np.mean(adaptive_scores):.2f}, Success={((adaptive_scores.count(0)/runs)*100:.0f}%)")
682
683 # Statistical test and visualization
684 scores_dict = {'Standard GA': standard_scores, 'Adaptive GA':
685 adaptive_scores}
686 compute_statistical_tests(scores_dict, output_file="adaptive_comparison_stats.txt")
687 plot_statistical_comparison(scores_dict,
688                               output_prefix="adaptive_comparison_stats",
689                               title=f"Standard vs Adaptive GA (N={num_nodes})")
690
691 # Plot convergence comparison (using median trace)
692 def get_median_trace(traces):
693     max_len = max(len(t) for t in traces)
694     padded = [t + [t[-1]] * (max_len - len(t)) for t in traces]
695     return np.median(padded, axis=0)
696
697 median_standard = get_median_trace(standard_traces)
698 median_adaptive = get_median_trace(adaptive_traces)
699
700 plt.figure(figsize=(10, 6))
701 plt.plot(median_standard, label='Standard GA (Fixed Mutation)',
```

```

699     linewidth=2)
700     plt.plot(median_adaptive, label='Adaptive GA (Diversity-Based)', linewidth=2, linestyle='--')
701     plt.title(f'GA Convergence: Standard vs Adaptive Mutation (N={num_nodes})')
702     plt.xlabel('Generation')
703     plt.ylabel('Broken Edges (Cost)')
704     plt.ylim(bottom=-0.5)
705     plt.legend()
706     plt.grid(True)
707
708     plt.savefig('experiment_5_adaptive_comparison.png', dpi=300,
709                 bbox_inches='tight')
710     print("Saved 'experiment_5_adaptive_comparison.png'")
711
712     # Save detailed results to file
713     with open('statistical_results.txt', 'a') as f:
714         f.write(f"\n{'*60}\n")
715         f.write(f"Adaptive Mutation Comparison (N={num_nodes}, p={prob})\n")
716         f.write(f"{'='*60}\n")
717         f.write(f"Standard GA: Mean={np.mean(standard_scores):.2f}, Success={(standard_scores.count(0)/runs)*100:.0f}%\n")
718         f.write(f"Adaptive GA: Mean={np.mean(adaptive_scores):.2f}, Success={(adaptive_scores.count(0)/runs)*100:.0f}%\n")
719     print("Results appended to 'statistical_results.txt'")
720
721 # =====
722 # MAIN EXECUTION
723 # =====
724
725 if __name__ == "__main__":
726     parser = argparse.ArgumentParser(description="Hamiltonian Path Metaheuristic Analysis")
727     parser.add_argument("-N", "--nodes", type=int, default=50, help="Number of nodes in the graph (default: 50)")
728     parser.add_argument("-p", "--prob", type=float, default=0.1, help="Edge creation probability (default: 0.1)")
729     parser.add_argument("--mode", type=str, default="all",
730                         choices=["all", "batch", "opt", "phase", "visual",
731                                  "stats", "adaptive"],
732                         help="Experiment mode to run individually")
733
734     args = parser.parse_args()
735
736     print(f"Running experiments with N={args.nodes} and p={args.prob}, Mode={args.mode}")
737
738     # 1. Main Baseline (using 'Swap')
739     if args.mode in ["all", "batch"]:
740         run_batch_experiments(args.nodes, args.prob)
741
742     # 2. The Report Recommendation (Proving 2-Opt is better)
743     if args.mode in ["all", "opt"]:
744         run_optimization_comparison(args.nodes)
745
746     # 3. Physics/Difficulty Analysis (Runs ALL algorithms across

```

```

    densities)
745   if args.mode in ["all", "phase"]:
746     run_phase_transition(args.nodes)
747
748 # 4. Visual Output
749 if args.mode in ["all", "visual"]:
750   save_best_graph_html(args.nodes, args.prob)
751
752 # 5. Statistical Analysis Mode (batch + stats)
753 if args.mode == "stats":
754   print("\n--- RUNNING BATCH WITH STATISTICAL ANALYSIS ---")
755   OPERATOR = "inversion"
756   g_exp = nx.erdos_renyi_graph(n=args.nodes, p=args.prob, seed=
757 SEED)
758
759   def run_batch_for_stats(algo_func, name):
760     scores = []
761     for _ in range(30):
762       _, cost, _ = algo_func(g_exp)
763       scores.append(cost)
764     return scores
765
766   sa_scores = run_batch_for_stats(lambda g:
767 run_simulated_annealing(g, max_steps=3000, operator=OPERATOR), "SA")
768   tabu_scores = run_batch_for_stats(lambda g: run_tabu_search(g,
769 max_steps=1500, operator=OPERATOR), "Tabu")
770   ga_scores = run_batch_for_stats(lambda g: run_genetic_algorithm(
771 g, generations=1500, operator=OPERATOR), "GA")
772
773   scores_dict = {'SA': sa_scores, 'Tabu': tabu_scores, 'GA':
774 ga_scores}
775   compute_statistical_tests(scores_dict)
776   plot_statistical_comparison(scores_dict,
777                               output_prefix=f"stats_N{args.nodes}_
778 {p}{args.prob}",
779                               title=f"Algorithm Comparison (N={
780 args.nodes}, p={args.prob})")
781
782 # 6. Adaptive Mutation Comparison
783 if args.mode == "adaptive":
784   run_adaptive_comparison(args.nodes, prob=args.prob)

```

Listing 7: Python 3.14.2