



Metaheuristic approach to the Hamiltonian Path

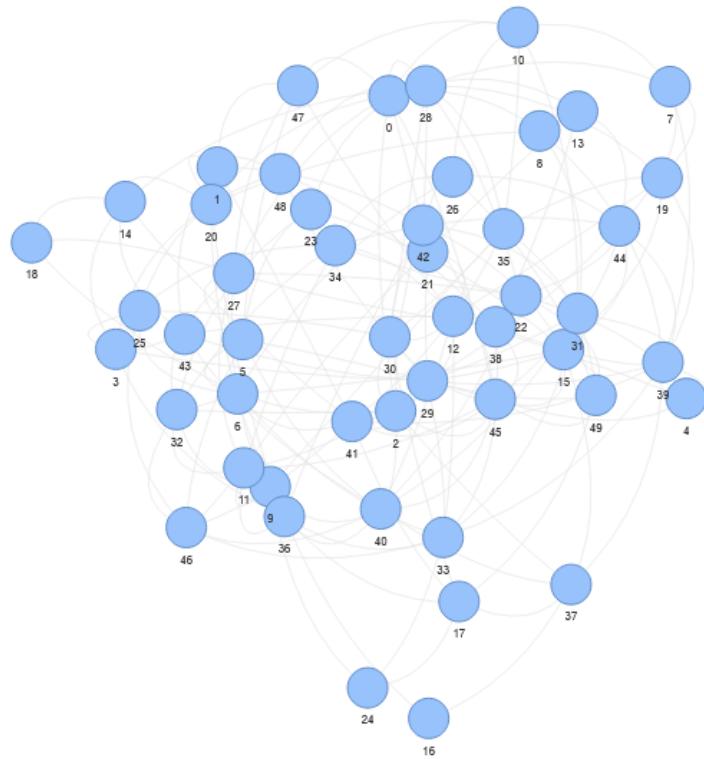


Figure 1: Visualised with Pyvis, depicts the best Hamiltonian path that could be solved for. The GA found using the 2-OPT (inversion) operator in the Erdős-Rényi random graph. Broken edges are depicted in red.

Module: Algorithms and Combinatorial Thinking 2025-2026

Written by:
Christos Christophoros Mitsakopoulos

Supervised by: Dr. Franck Delaplace

January 20, 2026

Contents

1	Formal Definition and Introduction	2
2	Random Generation of the Problem: Erdős-Rényi Random Graph	2
3	Test Environment: Experimental Choices and Reasoning	3
4	Metaheuristic Algorithms: Mode of Action and Code Snippets	4
4.1	Objective Function	4
4.2	Simulated Annealing (SA)	4
4.3	Tabu Search (TS)	4
4.4	Genetic Algorithm (GA)	5
4.4.1	Adaptive Mutation Rate	5
5	Experiment: Comparing the Metaheuristic Algorithms	6
5.1	Batch Testing	6
5.2	Phase Transition	7
6	GA: Comparing Swap Operator to 2-Opt	8
6.1	Statistical Validation	9
7	Scalability Analysis	10
8	Annex	10

1 Formal Definition and Introduction

This report evaluates three metaheuristic approaches: *Simulated Annealing*, *Tabu Search*, and a *Genetic Algorithm*.

To start with, the Hamiltonian Path problem asks if a graph $G = (V, E)$ contains a path that visits every vertex / node $v \in V$ exactly once.

Given that the Hamiltonian Path problem is NP-Complete, as the number of vertices of the vertex set $n = |V|$ increases, the computational complexity required to solve the problem via brute force grows factorially ($O(n!)$). Consequently, exact algorithms become computationally intractable for large n , necessitating the use of metaheuristics – like the ones tested in this report. For this exact reason no evaluation of a brute force approach has been considered in the experiments of this report, as the computational and time complexity required is too high to practically implement.

In the context of bioinformatics, specifically *de novo* genome assembly, this problem is interesting. Arranging DNA reads in an acyclic graph can be modelled as finding a Hamiltonian path that maximises read overlap to create a contiguous sequence. Since this is a bioinformatics course, a genetic algorithm appropriated from the course material (much like all other algorithms implemented) was tested to match the theme of genome assembly heuristics.

Find all relevant results / as well as reproduce the experiment using the `main.py` at the following address: https://github.com/cmitsakopoulos/Delaplace_coursework. The Python script automates all tests demonstrated in this report and moreover, accepts user arguments through the command line interface; intended for tweaking parameters regarding base graph generation. The `-mode` argument allows running specific experiments:

```
1 # Default: runs all standard experiments
2 python main.py --nodes 50 --prob 0.15
3
4 # Specific modes:
5 python main.py --mode batch      # Batch testing only
6 python main.py --mode stats     # Statistical tests (Kruskal-Wallis,
    Wilcoxon)
7 python main.py --mode adaptive # Compare standard vs adaptive mutation
    GA
```

Listing 1: Available experiment modes. The default mode "all" runs batch, optimisation, phase transition, and visualisation experiments.

2 Random Generation of the Problem: Erdős-Rényi Random Graph

To benchmark the metaheuristic algorithms, an Erdős-Rényi (ER) $G(n, p)$ model was used. In which:

- n : The number of vertices in the graph.
- p : The probability that an edge exists between any two distinct vertices.

A larger parameter p will by effect increase the likelihood of finding a Hamiltonian Path, given that the probability of any two nodes having a connecting edge is larger. This was evident when first trying out the `main.py` Python script, where preliminary tests with a p of ≈ 0.3 and $n = 50$ showed that a Hamiltonian Path was indeed mathematically possible to find. It was so much so that all algorithms would converge to a zero cost solution (no broken edges).

```
1 # n_default = 50, p_default = 0.1, seed hardcoded to 42
2 g_exp = nx.erdos_renyi_graph(n=num_nodes, p=prob, seed=SEED)
```

Listing 2: Small snippet of code in which NetworkX in Python can generate an ER graph of chosen parameters, identify that the default options are intended for a "challenging" benchmark for the algorithms. Hardcoding the seed was not a deliberate analytical choice, just to ensure reproducibility of the graph itself – the algorithms are stochastic too...

3 Test Environment: Experimental Choices and Reasoning

To ensure significance of the results, each metaheuristic algorithm was executed for $N = 30$ runs. Relying on the Central Limit Theorem, the sampling distribution of the mean approximates a normal distribution as $N \geq 30$, even if the underlying population distribution is non-Gaussian. While increasing N reduces the Standard Error of the Mean (SEM), the precision improves only with the square root of N (i.e., $SEM = \sigma/\sqrt{N}$).

Additionally, in order to observe the impact on phase transition across all metaheuristics tested, the Komlós–Szemerédi theorem/threshold was used to compute the theoretical p -limit of the problem graph, specifically using the following equation:

$$p_{\text{limit}} = \frac{\ln(n) + \ln(\ln(n))}{n} \quad (1)$$

While $p < p_{\text{limit}}$, the probability of a Hamiltonian path existing approaches zero; thus any algorithm claiming a valid solution is likely "hallucinating", being overconfident (due to verification errors), or has encountered a statistical anomaly (highly unlikely in my testing case). Conversely, for $p > p_{\text{limit}}$, a path almost surely exists: consequently, if an algorithm fails to converge to a zero-cost solution in this region (zero broken edges), it shows that the algorithm is underperforming rather than the problem being impossible.

Another important consideration, was to examine the impact of switching from a "Swap" operator to an **"Inversion" (2-OPT) operator** between testing cases. This was also tested with the GA to demonstrate how 2-OPT can be a positive addition to an already accurate algorithm. Additionally, it was tested with SA and TS to demonstrate that the impact of this change is not algorithm-specific. This phenomenon should become evident when looking at the **"Phase Transition"** experiment of this report, where the 2-OPT operator was used instead of the Swap operator.

To move beyond visual comparisons and quantify whether observed performance differences are statistically meaningful, non-parametric statistical tests were employed. The Kruskal-Wallis H-test determines whether at least one algorithm performs significantly

differently from the others, while pairwise Wilcoxon signed-rank tests identify which specific pairs differ. These non-parametric tests are appropriate because solution quality scores do not follow a normal distribution—algorithms either find optimal solutions (cost = 0) or get stuck at various local optima. Additionally, Cohen’s d effect size quantifies the practical magnitude of differences, distinguishing between statistically significant results and practically meaningful ones. An effect size $|d| \geq 0.8$ indicates a large, practically important difference between algorithms.

4 Metaheuristic Algorithms: Mode of Action and Code Snippets

4.1 Objective Function

The aim of these algorithms is to identify a permutation S of vertices that minimises the number of broken edges in the path. For a graph $G = (V, E)$ and a candidate path $S = [v_1, v_2, \dots, v_n]$, the cost function $C(S)$ is defined as:

$$C(S) = (n - 1) - \sum_{i=1}^{n-1} \mathbb{I}((v_i, v_{i+1}) \in E) \quad (2)$$

Where \mathbb{I} is an indicator function that equals 1 if the edge exists and 0 otherwise. A global optimum is reached when $C(S) = 0$ – a Hamiltonian Path.

4.2 Simulated Annealing (SA)

Simulated Annealing explores solutions by accepting both better or even worse solutions, based on a probability that decreases over time (referred to as parameter T). This logic prevents the algorithm from arriving at a final solution before reaching a truly optimal or near optimal solution. The probability P of accepting a new solution S' with cost difference $\Delta C = C(S') - C(S)$, is governed by the Metropolis criterion (see also the Python implementation below):

$$P(\text{accept}) = \begin{cases} 1 & \text{if } \Delta C < 0 \\ e^{-\frac{\Delta C}{T}} & \text{if } \Delta C \geq 0 \end{cases} \quad (3)$$

```

1 # From main.py: Calculate cost difference
2 delta = neighbor_cost - current_cost
3
4 # Accept if better (delta < 0) OR with probability exp(-delta/T)
5 if delta < 0 or random.random() < exp(-delta / temp):
6     current = neighbor
7     current_cost = neighbor_cost

```

Listing 3: *Metropolis Criterion: As the loop progresses, the temperature T decays geometrically ($T_{k+1} = 0.985 \cdot T_k$), gradually turning the search into a simple greedy descent (hill climbing)*

4.3 Tabu Search (TS)

Tabu Search differs from SA by using a deterministic, memory-based approach. TS explores the immediate neighbourhood of a current solution, then moves to the best available neighbouring solution, even if that neighbouring solution is worse than the current solution.

To prevent cycling (revisiting the same solutions again and again), the algorithm maintains a *Tabu List* – a short-term memory that prevents recent solutions for a specific duration, called tenure.

```
1 # From main.py: Moving to the best candidate not in the Tabu list
2 if cand not in tabu_list:
3     current = cand
4     found_move = True
5
6 # Update memory
7 tabu_list.append(current)
8 if len(tabu_list) > tenure:
9     tabu_list.pop(0) # Remove oldest entry
```

Listing 4: Tabu Search Memory Logic

4.4 Genetic Algorithm (GA)

The Genetic Algorithm attempts to mimic natural selection. Unlike SA and TS, which improve a single solution, GA evolves a population of solutions. The operator for permutation is *Ordered Crossover* (OX1), which is important because standard single-point crossover would result in duplicate or missing vertices / nodes.

The code uses OX1 to preserve the ordering of a sub-segment from one "parent" while filling the remaining slots with genes from the second "parent".

```
1 # From main.py: Preserves sub-segment from parent 1
2 child[start:end] = p1[start:end]
3
4 # Fills remaining slots from parent 2, skipping duplicates
5 for i in range(size):
6     if child[i] is None:
7         while p2[current_p2_idx] in child:
8             current_p2_idx += 1
9         child[i] = p2[current_p2_idx]
```

Listing 5: Ordered Crossover (OX1) Implementation

Using the standard `swap` operator (exchanging two indices) disrupts the adjacency of the path significantly. In comparison, the `inversion` (2-Opt) operator reverses a segment of the path. This is mathematically better for path search problems because it preserves the internal adjacency of the reversed segment, only breaking the two edges at the endpoints of the segment. This distinction is implemented via the ‘`op_inversion`’ function in ‘`main.py`’ and is the primary driver for convergence in denser graphs. A graph is produced at the end during testing to demonstrate the differences of both applications.

$$\text{Swap}(S) \rightarrow \text{High disruption of edges} \quad (4)$$

$$\text{Inversion}(S) \rightarrow \text{Minimal disruption (2-Opt)} \quad (5)$$

4.4.1 Adaptive Mutation Rate

Another improvement over that of the mutation operator (Swap vs 2-Opt), which determines how to mutate a solution in order to develop it, is improving on how often to mutate. With a high fixed mutation rate, the algorithm promotes exploration but disrupts good solutions, while low mutation rates enable exploitation but risk converging early; thereby leading to suboptimal solutions. An adaptive approach addresses this by adjusting the mutation probability μ based on population diversity D :

$$D = \frac{|\mathcal{F}_{\text{unique}}|}{P}, \quad \mu_{\text{adaptive}} = \mu_{\text{base}} \cdot (1 + (0.5 - D)) \quad (6)$$

When $D < 0.5$ (population becoming homogeneous), mutation increases to reintroduce variation. When $D > 0.5$ (healthy diversity), mutation decreases to exploit promising solutions. The rate is clamped to $[0.1, 0.6]$. This self-regulating mechanism helps prevent premature convergence without manual tuning.

```

1 def calculate_diversity(fitnesses):
2     return len(set(fitnesses)) / len(fitnesses)
3
4 # Inside GA loop, per generation:
5 if adaptive:
6     diversity = calculate_diversity(fitnesses)
7     current_mut = base_mut_rate * (1 + (0.5 - diversity))
8     current_mut = max(0.1, min(0.6, current_mut))

```

Listing 6: *Adaptive mutation rate logic from main.py. Diversity is computed per generation, and mutation rate is adjusted accordingly.*

5 Experiment: Comparing the Metaheuristic Algorithms

5.1 Batch Testing

The experiment code in `main.py` has been hardcoded to repeat the stochastic run of each metaheuristic algorithm 30 times, with the user chosen parameters upon initialization through the command line. The results depicted in [Figure 2](#) and [Figure 3](#), were computed with an ER problem graph of $n = 50$ and $p = 0.15$.

Briefly, with reference to [Figure 2](#), the lowest solution quality (highest number of broken edges) is demonstrated by SA, then TS and with the GA, being the best performer. Inversely, the highest execution time per batch run, was demonstrated by GA, followed by TS and lastly SA. These results confirm the expected correlation between computational cost and solution quality (indicating the code works as intended). SA is the fastest because it performs a single $O(1)$ evaluation per step, but its single-trajectory stochasticity struggles to escape bad solutions (local optima) in a solution landscape. TS evaluates a neighbourhood of size $k = 50$ per step ($O(k)$) (focus on local searches), which linearly increases runtime. Lastly, the GA achieves a near-optimal solution albeit, at the highest computational cost. It maintains diversity through a population-based search ($O(P \cdot N)$ per generation), allowing it to traverse the complex fitness landscape more effectively than the trajectory-based methods.

Nevertheless, the results in [Figure 2](#) clearly show that all algorithms – apart from the GA – are greatly underperforming in this testing case. For an ER graph with $n = 50$, the

Komlós–Szemerédi limit is $p = 0.106$, considering that the chosen – testing – $p = 0.15$, is greater than the theoretical limit, a solution should be mathematically infeasible; albeit, difficult to obtain.




Figure 2: In the box whisker plots depicted, "Solution Quality" (Y-axis: Broken Edges) is shown on the left and "Execution Time" (Y-axis: Time(seconds)) on the right. As discussed prior, due to the low probability of edges between any two nodes, the graph is sparse and the metaheuristic algorithms – while performing differently from one another – are not successful in this testing case. As such, the amount of broken edges seen between algorithms is not unlikely, and the inverse relationship between the number of broken edges and execution time is a clear demonstration of algorithmic complexity – which leads to positive outcomes in problem solving.

5.2 Phase Transition

Briefly, with reference to Figure 3, there appears to be overconfidence of the TS and GA algorithms, which demonstrate a non-zero level of success at a mathematically improbable rate of $p = 0.1$; where $p = 0.1 < p_{\text{limit}} = 0.106$. Given that the difference between p and p_{limit} is $\Delta p = 0.006$, this could be believed to be a statistical anomaly that occurred in this new 30-run test – compared to the previous batch testing scenario. After introducing a new `argparse` argument to the `main.py` file to exclusively re-run the phase transition experiment, the results – over 5 re-runs – led to the same conclusions as in Figure 3. One could argue that due to a **switch from the "Swap" operator to the more optimised "Inversion" (2-Opt) operator**, the GA and TS algorithms are able to find a global optimum at a mathematically improbable rate of $p = 0.1$, as well as be considerably successful at higher graph densities ($p > 0.1$); all the while SA is equally performing better than before (than in Figure 2).

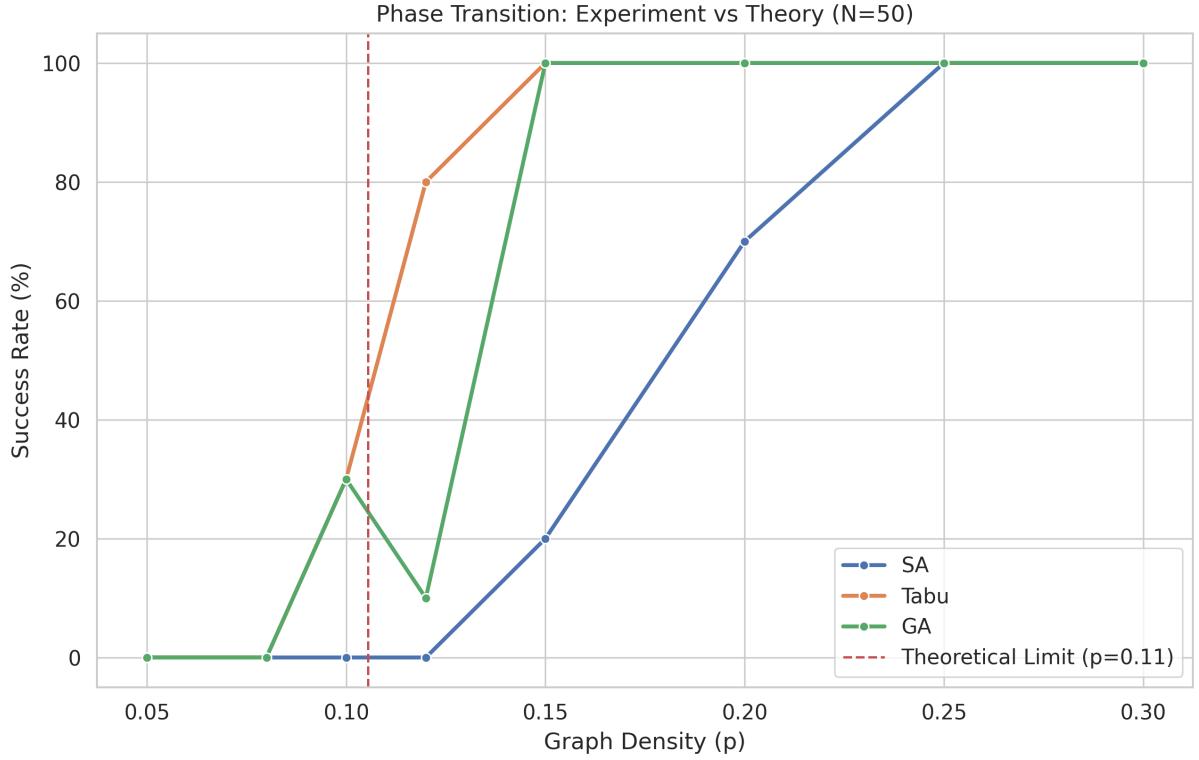


Figure 3: This line chart depicts the phase transition of the proposed problem Erdos-Renyi graph, at a cycling p rate between 0.05 and 0.3. The theoretical limit is computed automatically based on the input problem graph, given the Komlós–Szemerédi formula. Observe overconfidence from the TS and GA algorithms, which demonstrate success at a mathematically improbable rate of p . The success rate of all algorithms is congruent with previous results.

6 GA: Comparing Swap Operator to 2-Opt

Congruent with the phase transition results, it is evident that the switch from the Swap operator (Figure 2) to the Inversion (2-Opt) operator (Figure 3), improves the performance of the GA (see Figure 4). Importantly, looking at the number of broken edges at each iteration, it is clear that the 2-OPT based GA converges at a faster and smoother, or more consistent pace, than the Swap based GA. To briefly reiterate on what was said prior on the report, this phenomenon can be due to two complementary reasons:

Firstly, the 2-OPT (or inversion) operator is less destructive at each iteration. It reverses a segment between indices i and j at each iteration, removing two edges (v_i, v_{i+1}) and (v_j, v_{j+1}) , and replacing them with (v_i, v_j) and (v_{i+1}, v_{j+1}) , whilst preserving the adjacency of all nodes within the segment. As such, the 2-OPT operator enables the algorithm to make small changes which are beneficial at each iteration. Therefore, as can be seen in Figure 4, the GA with the 2-OPT operator improves its global solution at each iteration, whereas the Swap operator will repeatedly plateau/stagnate due to destructive solutions between iterations. Secondly, as mentioned before, the Swap operator exchanges two nodes v_i and v_j arbitrarily, which can affect up to four edges: (v_{i-1}, v_i) , (v_i, v_{i+1}) , (v_{j-1}, v_j) , and (v_j, v_{j+1}) . Instead of making minor improvements locally, the Swap operator will lead the algorithm to more destructive changes at each iteration. In turn, this will cause the search

to stagnate (maintain one solution over multiple iterations) over a certain amount of iterations. While a more beneficial solution can be found accidentally, the algorithm risks getting stuck in local minima; minima which can be avoided by the 2-OPT operator.

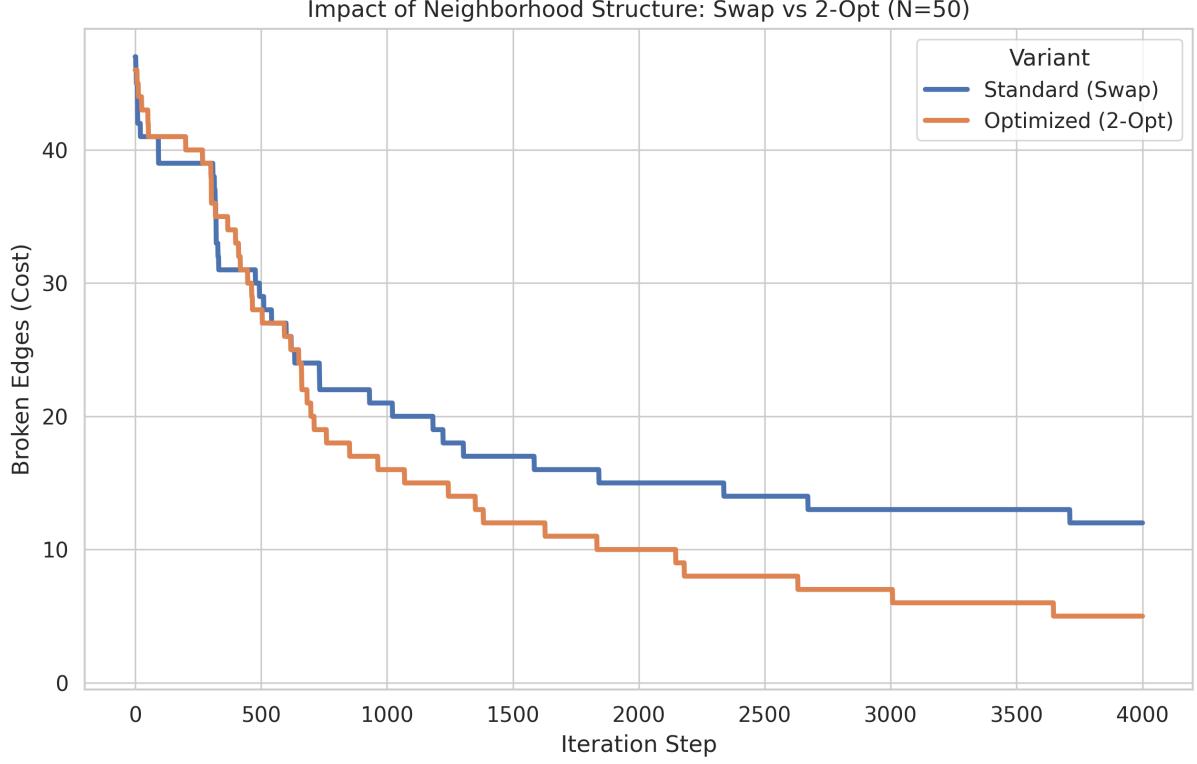


Figure 4: In the depicted line chart, Broken Edges (Y-axis) for which the objective function seeks to minimise (Cost) by maximising the number of edges in the path, is plotted against the number of generations – repeated runs (X-axis). The GA with the Swap operator is shown in blue, while the GA with the Inversion (2-Opt) operator is in orange. Observe the clear disparity between the two operators, in which it's clear that the 2-OPT operator GA converges much faster to a near-global optimum (zero broken edges) than its Swap-GA counterpart.

6.1 Statistical Validation

To quantify the statistical significance of the performance differences observed in the figures above, the Kruskal-Wallis H-test was applied:

$$H = \frac{12}{N(N+1)} \sum_{i=1}^k \frac{R_i^2}{n_i} - 3(N+1) \quad (7)$$

A significant result ($p < 0.05$) confirms that the algorithms perform differently. Pairwise Wilcoxon signed-rank tests further identify which specific pairs differ:

$$W = \min \left(\sum_{d_i > 0} R_i, \sum_{d_i < 0} R_i \right) \quad (8)$$

Finally, Cohen's d effect size quantifies the magnitude of these differences:

$$d = \frac{\bar{x}_1 - \bar{x}_2}{s_{\text{pooled}}} \quad \text{where} \quad s_{\text{pooled}} = \sqrt{\frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2}} \quad (9)$$

The statistical analysis is visualised through a three-panel figure containing: a box plot with Mann-Whitney significance annotations (star notation: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$), mean broken edges with 95% confidence interval error bars, and an effect size heatmap showing pairwise Cohen's d values.

Uncomment when figure is generated:




Figure 5: *Statistical comparison of SA, Tabu Search, and GA. Left: Box plot with Mann-Whitney significance annotations. Centre: Mean broken edges with 95% CI error bars. Right: Cohen's d effect size heatmap showing pairwise comparisons.*

7 Annex

```

1 import random
2 import time
3 import math
4 import itertools
5 import argparse
6 import numpy as np
7 import pandas as pd
8 import networkx as nx
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11 from math import exp, log
12
13 # --- CONFIGURATION ---
14 SNS_THEME = "whitegrid"
15 SNS_CONTEXT = "paper"
16 SEED = 42
17
18 # Apply settings
19 random.seed(SEED)
20 np.random.seed(SEED)
21 sns.set_theme(style=SNS_THEME, context=SNS_CONTEXT, font_scale=1.2)
22
23 # =====
24 # 1. CORE HELPER & FITNESS FUNCTIONS
25 # =====
26

```

```

27 def get_initial_solution(nodes):
28     """Returns a random permutation of nodes."""
29     p = list(nodes).copy()
30     random.shuffle(p)
31     return p
32
33 def count_edges(path, graph):
34     """Counts valid edges in the path for a specific graph."""
35     edges = 0
36     # Optimization: Iterate via index to avoid creating new lists
37     for i in range(len(path) - 1):
38         if graph.has_edge(path[i], path[i+1]):
39             edges += 1
40     return edges
41
42 def fitness_min(path, graph):
43     """Minimisation: Returns number of broken edges. Target = 0."""
44     target = len(path) - 1
45     valid = count_edges(path, graph)
46     return float(target - valid)
47
48 def fitness_max(path, graph):
49     """Maximisation: Returns number of valid edges. Target = N-1."""
50     return float(count_edges(path, graph))
51
52 # =====
53 # 2. OPERATORS (THE OPTIMISATIONS)
54 # =====
55
56 def op_swap(genome):
57     """Standard Swap: Exchanges two random nodes."""
58     n = len(genome)
59     i, j = random.sample(range(n), 2)
60     genome[i], genome[j] = genome[j], genome[i]
61     return genome
62
63 def op_inversion(genome):
64     """
65     2-Opt Inversion: Reverses a random sub-segment.
66     OPTIMIZATION: Preserves adjacency better than swap.
67     """
68     n = len(genome)
69     i, j = sorted(random.sample(range(n), 2))
70     # Reverse the segment between i and j
71     genome[i:j+1] = genome[i:j+1][::-1]
72     return genome
73
74 # =====
75 # 3. METAHEURISTIC ALGORITHMS
76 # =====
77
78 def run_simulated_annealing(graph, max_steps=3000, temp0=100.0, operator="swap"):
79     """Simulated Annealing with selectable operator."""
80     nodes = list(graph.nodes)
81     current = get_initial_solution(nodes)
82     best = current.copy()
83

```

```

84     current_cost = fitness_min(current, graph)
85     best_cost = current_cost
86     trace = [best_cost]
87     temp = temp0
88
89     mutate_func = op_inversion if operator == "inversion" else op_swap
90
91     for step in range(max_steps):
92         if best_cost == 0: break
93
94         neighbor = current.copy()
95         neighbor = mutate_func(neighbor)
96
97         neighbor_cost = fitness_min(neighbor, graph)
98         delta = neighbor_cost - current_cost
99
100        if delta < 0 or random.random() < exp(-delta / temp):
101            current = neighbor
102            current_cost = neighbor_cost
103            if current_cost < best_cost:
104                best = current.copy()
105                best_cost = current_cost
106
107            trace.append(best_cost)
108            temp *= 0.985
109
110    return best, best_cost, trace
111
112 def run_tabu_search(graph, max_steps=1000, tenure=20, operator="swap"):
113     """Tabu Search with selectable operator."""
114     nodes = list(graph.nodes)
115     current = get_initial_solution(nodes)
116     best = current.copy()
117     best_cost = fitness_min(best, graph)
118
119     tabu_list = []
120     trace = [best_cost]
121     mutate_func = op_inversion if operator == "inversion" else op_swap
122
123     for step in range(max_steps):
124         if best_cost == 0: break
125
126         candidates = []
127         for _ in range(50):
128             cand = current.copy()
129             cand = mutate_func(cand)
130             candidates.append(cand)
131
132         candidates.sort(key=lambda p: fitness_min(p, graph))
133
134         found_move = False
135         for cand in candidates:
136             cand_cost = fitness_min(cand, graph)
137             if cand_cost < best_cost:
138                 current = cand
139                 best = cand
140                 best_cost = cand_cost
141                 found_move = True

```

```

142         break
143     if cand not in tabu_list:
144         current = cand
145         found_move = True
146         break
147
148     trace.append(best_cost)
149     if found_move:
150         tabu_list.append(current)
151         if len(tabu_list) > tenure:
152             tabu_list.pop(0)
153
154     return best, best_cost, trace
155
156 def run_genetic_algorithm(graph, pop_size=200, generations=1000,
157 cross_rate=0.8, mut_rate=0.3, operator="swap"):
158     """Genetic Algorithm with selectable operator."""
159     nodes = list(graph.nodes)
160     target = len(nodes) - 1
161
162     mutate_func = op_inversion if operator == "inversion" else op_swap
163
164     def ordered_crossover(p1, p2):
165         size = len(p1)
166         start, end = sorted(random.sample(range(size), 2))
167         child = [None] * size
168         child[start:end] = p1[start:end]
169         current_p2_idx = 0
170         for i in range(size):
171             if child[i] is None:
172                 while p2[current_p2_idx] in child:
173                     current_p2_idx += 1
174                 child[i] = p2[current_p2_idx]
175
176     population = [get_initial_solution(nodes) for _ in range(pop_size)]
177     best_sol = None
178     best_broken = float('inf')
179     trace = []
180
181     for gen in range(generations):
182         fitnesses = [fitness_max(p, graph) for p in population]
183         best_val = max(fitnesses)
184         current_broken = target - best_val
185
186         if current_broken < best_broken:
187             best_broken = current_broken
188             best_sol = population[fitnesses.index(best_val)]
189
190         trace.append(best_broken)
191         if current_broken == 0: break
192
193         next_pop = [best_sol.copy()]
194         while len(next_pop) < pop_size:
195             competitors = random.sample(population, 3)
196             winner = max(competitors, key=lambda p: fitness_max(p, graph))
197             next_pop.append(winner.copy())

```

```

198     population = next_pop
199
200     for i in range(1, pop_size - 1, 2):
201         if random.random() < cross_rate:
202             population[i] = ordered_crossover(population[i],
203                                                 population[i+1])
204
205         for i in range(1, pop_size):
206             if random.random() < mut_rate:
207                 population[i] = mutate_func(population[i])
208
209     return best_sol, best_broken, trace
210
211 # =====
212 # 4. EXPERIMENT RUNNERS
213 # =====
214
215 def run_batch_experiments(num_nodes, prob):
216     """Runs each algorithm 30 times and plots results."""
217     # Baseline comparison using Standard Swap
218     OPERATOR = "swap"
219
220     g_exp = nx.erdos_renyi_graph(n=num_nodes, p=prob, seed=SEED)
221
222     print(f"\n--- 1. BATCH EXPERIMENT (Baseline 'Swap', N={num_nodes}, p={prob}, 30 Runs) ---")
223
224     def run_batch(algo_func, name):
225         scores = []
226         times = []
227         successes = 0
228         for _ in range(30):
229             start = time.time()
230             _, cost, _ = algo_func(g_exp)
231             dur = time.time() - start
232             scores.append(cost)
233             times.append(dur)
234             if cost == 0: successes += 1
235
236         rate = (successes/30)*100
237         print(f"{name}: Mean Cost={np.mean(scores):.2f}, Success={rate:.1f}%, Mean Time={np.mean(times):.4f}s")
238         return scores, times
239
240     sa_scores, sa_times = run_batch(lambda g: run_simulated_annealing(g,
241                                                                           max_steps=3000, operator=OPERATOR), "SA")
242     tabu_scores, tabu_times = run_batch(lambda g: run_tabu_search(g,
243                                                                           max_steps=1500, operator=OPERATOR), "Tabu")
244     ga_scores, ga_times = run_batch(lambda g: run_genetic_algorithm(g,
245                                                                           generations=1500, operator=OPERATOR), "GA")
246
247     data = {
248         'Algorithm': ['SA']*30 + ['Tabu']*30 + ['GA']*30,
249         'Broken Edges': sa_scores + tabu_scores + ga_scores,
250         'Time (s)': sa_times + tabu_times + ga_times
251     }
252     df = pd.DataFrame(data)

```

```

250     fig, axes = plt.subplots(1, 2, figsize=(14, 6))
251
252     sns.boxplot(data=df, x='Algorithm', y='Broken Edges', hue='Algorithm',
253                 legend=False, ax=axes[0], palette="viridis")
254     axes[0].set_title(f'Solution Quality (N={num_nodes}, Lower is Better)')
255     axes[0].set_ylim(bottom=-0.5)
256
257     sns.boxplot(data=df, x='Algorithm', y='Time (s)', hue='Algorithm',
258                 legend=False, ax=axes[1], palette="magma")
259     axes[1].set_title(f'Execution Time (N={num_nodes})')
260     axes[1].set_ylim(bottom=0)
261
262     plt.tight_layout()
263     plt.savefig('experiment_1_batch_results.png', dpi=300, bbox_inches='tight')
264     print("Saved 'experiment_1_batch_results.png'")
265     # plt.show() # Uncomment if running interactively
266
267 def run_optimization_comparison(num_nodes):
268     """Compares Standard 'Swap' vs Optimized 'Inversion' (2-Opt)."""
269     print(f"\n--- 2. OPTIMISATION EVALUATION (N={num_nodes}) ---")
270     g_opt = nx.erdos_renyi_graph(n=num_nodes, p=0.1, seed=SEED)
271
272     print("Running SA with Standard Swap...")
273     _, _, trace_swap = run_simulated_annealing(g_opt, max_steps=4000,
274                                                 operator="swap")
275
276     print("Running SA with Optimized Inversion (2-Opt)...")
277     _, _, trace_inv = run_simulated_annealing(g_opt, max_steps=4000,
278                                                operator="inversion")
279
280     df_opt = pd.DataFrame({
281         'Step': list(range(len(trace_swap))) + list(range(len(trace_inv))),
282         'Broken Edges': trace_swap + trace_inv,
283         'Variant': ['Standard (Swap)'] * len(trace_swap) + ['Optimized (2-Opt)'] * len(trace_inv)
284     })
285
286     plt.figure(figsize=(10, 6))
287     sns.lineplot(data=df_opt, x='Step', y='Broken Edges', hue='Variant',
288                  linewidth=2.5)
289     plt.title(f'Impact of Neighborhood Structure: Swap vs 2-Opt (N={num_nodes})')
290     plt.ylabel('Broken Edges (Cost)')
291     plt.xlabel('Iteration Step')
292     plt.ylim(bottom=-0.5)
293
294     plt.savefig('experiment_2_optimization_comparison.png', dpi=300,
295     bbox_inches='tight')
296     print("Saved 'experiment_2_optimization_comparison.png'")
297
298 def run_phase_transition(num_nodes):
299     """
300     Analyzes difficulty vs graph density across ALL algorithms.
301     This provides a comprehensive view of algorithmic limits.
302 
```

```

298 """
299 print(f"\n--- 3. PHASE TRANSITION ANALYSIS (N={num_nodes}) ---")
300 # Densities to test
301 densities = [0.05, 0.08, 0.1, 0.12, 0.15, 0.2, 0.25, 0.3]
302
303 # Calculate Theoretical Threshold for THIS N
304 # Koml s & Szemer di theorem: p = (ln(n) + ln(ln(n))) / n
305 if num_nodes > 1:
306     math_threshold = (log(num_nodes) + log(log(num_nodes))) /
307     num_nodes
308     print(f"Theoretical Critical Threshold for N={num_nodes}: p ~ {math_threshold:.3f}")
309 else:
310     math_threshold = 0
311
312 # Store results for plotting
313 results = {'Density': [], 'Success Rate': [], 'Algorithm': []}
314
315 # Test all algorithms to see which one handles the transition best
316 algorithms = {
317     'SA': lambda g: run_simulated_annealing(g, max_steps=2500,
318 operator="inversion"),
319     'Tabu': lambda g: run_tabu_search(g, max_steps=1500, operator="inversion"),
320     'GA': lambda g: run_genetic_algorithm(g, generations=1000,
321 operator="inversion")
322 }
323
324 for p in densities:
325     # Create a graph for this density
326     # Note: Ideally we average over multiple graphs, but for speed
327     # we use one seed per density
328     # or we generate a fresh one each run. Let's stick to one graph
329     # instance per density
330     # but 10 runs per algorithm on it.
331     g = nx.erdos_renyi_graph(n=num_nodes, p=p, seed=999)
332
333     print(f"Testing Density p={p}...")
334
335     for name, algo_func in algorithms.items():
336         successes = 0
337         for _ in range(10): # 10 runs per density/algorithm pair
338             _, cost, _ = algo_func(g)
339             if cost == 0: successes += 1
340
341         rate = (successes/10)*100
342         results['Density'].append(p)
343         results['Success Rate'].append(rate)
344         results['Algorithm'].append(name)
345
346 df_phase = pd.DataFrame(results)
347
348 plt.figure(figsize=(10, 6))
349 sns.lineplot(data=df_phase, x='Density', y='Success Rate', hue='Algorithm',
350 marker='o', linewidth=2)
351
352 # Plot the Theoretical Line
353 plt.axvline(x=math_threshold, color='r', linestyle='--', label=f'
```

```

Theoretical Limit (p={math_threshold:.2f}))')

348
349 plt.title(f'Phase Transition: Experiment vs Theory (N={num_nodes})')
350 plt.xlabel('Graph Density (p)')
351 plt.ylabel('Success Rate (%)')
352 plt.ylim(-5, 105)
353 plt.legend()
354 plt.grid(True)
355
356 plt.savefig('experiment_3_phase_transition.png', dpi=300,
bbox_inches='tight')
357 print("Saved 'experiment_3_phase_transition.png'")

358
359 def save_best_graph_html(num_nodes, prob):
    """Runs GA once and saves the result to HTML."""
360    print("\n--- 4. GENERATING VISUALISATION ---")
361    try:
362        from pyvis.network import Network
363    except ImportError:
364        print("Pyvis not installed. Skipping.")
365        return
366
367    # Use arguments, but if probability is too low, we might not find a
368    # path
369    g = nx.erdos_renyi_graph(n=num_nodes, p=prob, seed=SEED)
370    path, cost, _ = run_genetic_algorithm(g, generations=2000, operator=
371                                         "inversion")
372
373    print(f"Final Path Cost: {cost}")
374
375    net = Network(height="600px", width="100%", cdn_resources='remote')
376
377    for n in g.nodes:
378        net.add_node(int(n), label=str(n), color='#97c2fc')
379
380    for u, v in g.edges:
381        net.add_edge(int(u), int(v), color="#e0e0e0", width=1)
382
383    for i in range(len(path) - 1):
384        u, v = int(path[i]), int(path[i+1])
385        if g.has_edge(u, v):
386            net.add_edge(u, v, color='red', width=4)
387        else:
388            net.add_edge(u, v, color='red', width=4, dashes=True)
389
390    net.show("hamiltonian_path.html", notebook=False)
391    print("Saved to 'hamiltonian_path.html'")

392 # =====
393 # MAIN EXECUTION
394 # =====

395
396 if __name__ == "__main__":
397     parser = argparse.ArgumentParser(description="Hamiltonian Path
398     Metaheuristic Analysis")
399     parser.add_argument("-N", "--nodes", type=int, default=50, help="Number
      of nodes in the graph (default: 50)")
      parser.add_argument("-p", "--prob", type=float, default=0.1, help="
```

```

400     Edge creation probability (default: 0.1)")
401     parser.add_argument("--mode", type=str, default="all", choices=["all",
402                         "batch", "opt", "phase", "visual"], help="Experiment mode to run
403                         individually")
404
405     args = parser.parse_args()
406
407     print(f"Running experiments with N={args.nodes} and p={args.prob},
408           Mode={args.mode}")
409
410     # 1. Main Baseline (using 'Swap')
411     if args.mode in ["all", "batch"]:
412         run_batch_experiments(args.nodes, args.prob)
413
414     # 2. The Report Recommendation (Proving 2-Opt is better)
415     if args.mode in ["all", "opt"]:
416         run_optimization_comparison(args.nodes)
417
418     # 3. Physics/Difficulty Analysis (Runs ALL algorithms across
419     # densities)
420     if args.mode in ["all", "phase"]:
421         run_phase_transition(args.nodes)
422
423     # 4. Visual Output
424     if args.mode in ["all", "visual"]:
425         save_best_graph_html(args.nodes, args.prob)

```

Listing 7: Python 3.14.2